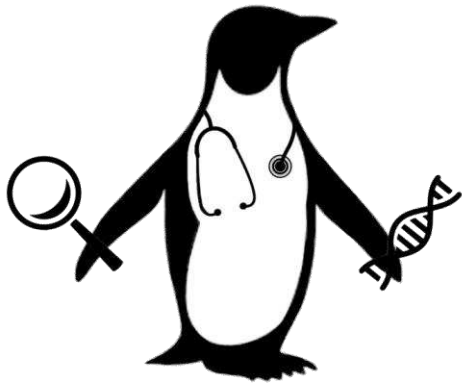

The PENGQUIN Odyssey:

A Quest for Query-able Semantic Genomic Data for Clinical Practice

Elias Crum

SIB Vital-IT Group Meeting
April 4th, 2025

SIB Research Stay: 1 March – 30 April

A Quick Thank You!

Tarcisio Mendes

Ana-Claudia Sima

Jerven Bolleman

Marco Pagni

An Intro to Elias

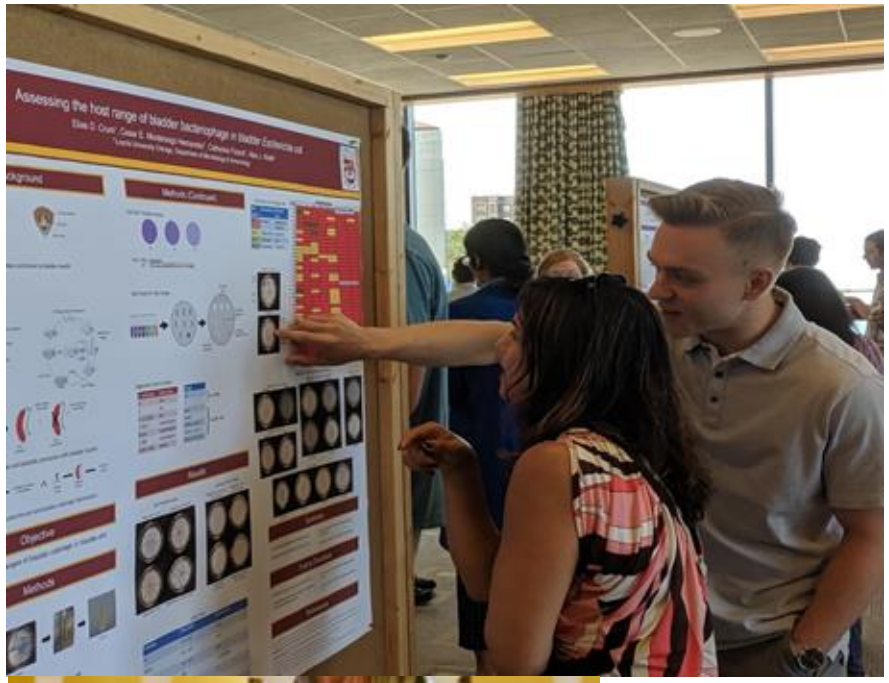

My Journey:

BSc + MSc - Bioinformatics

USA Geography Lesson

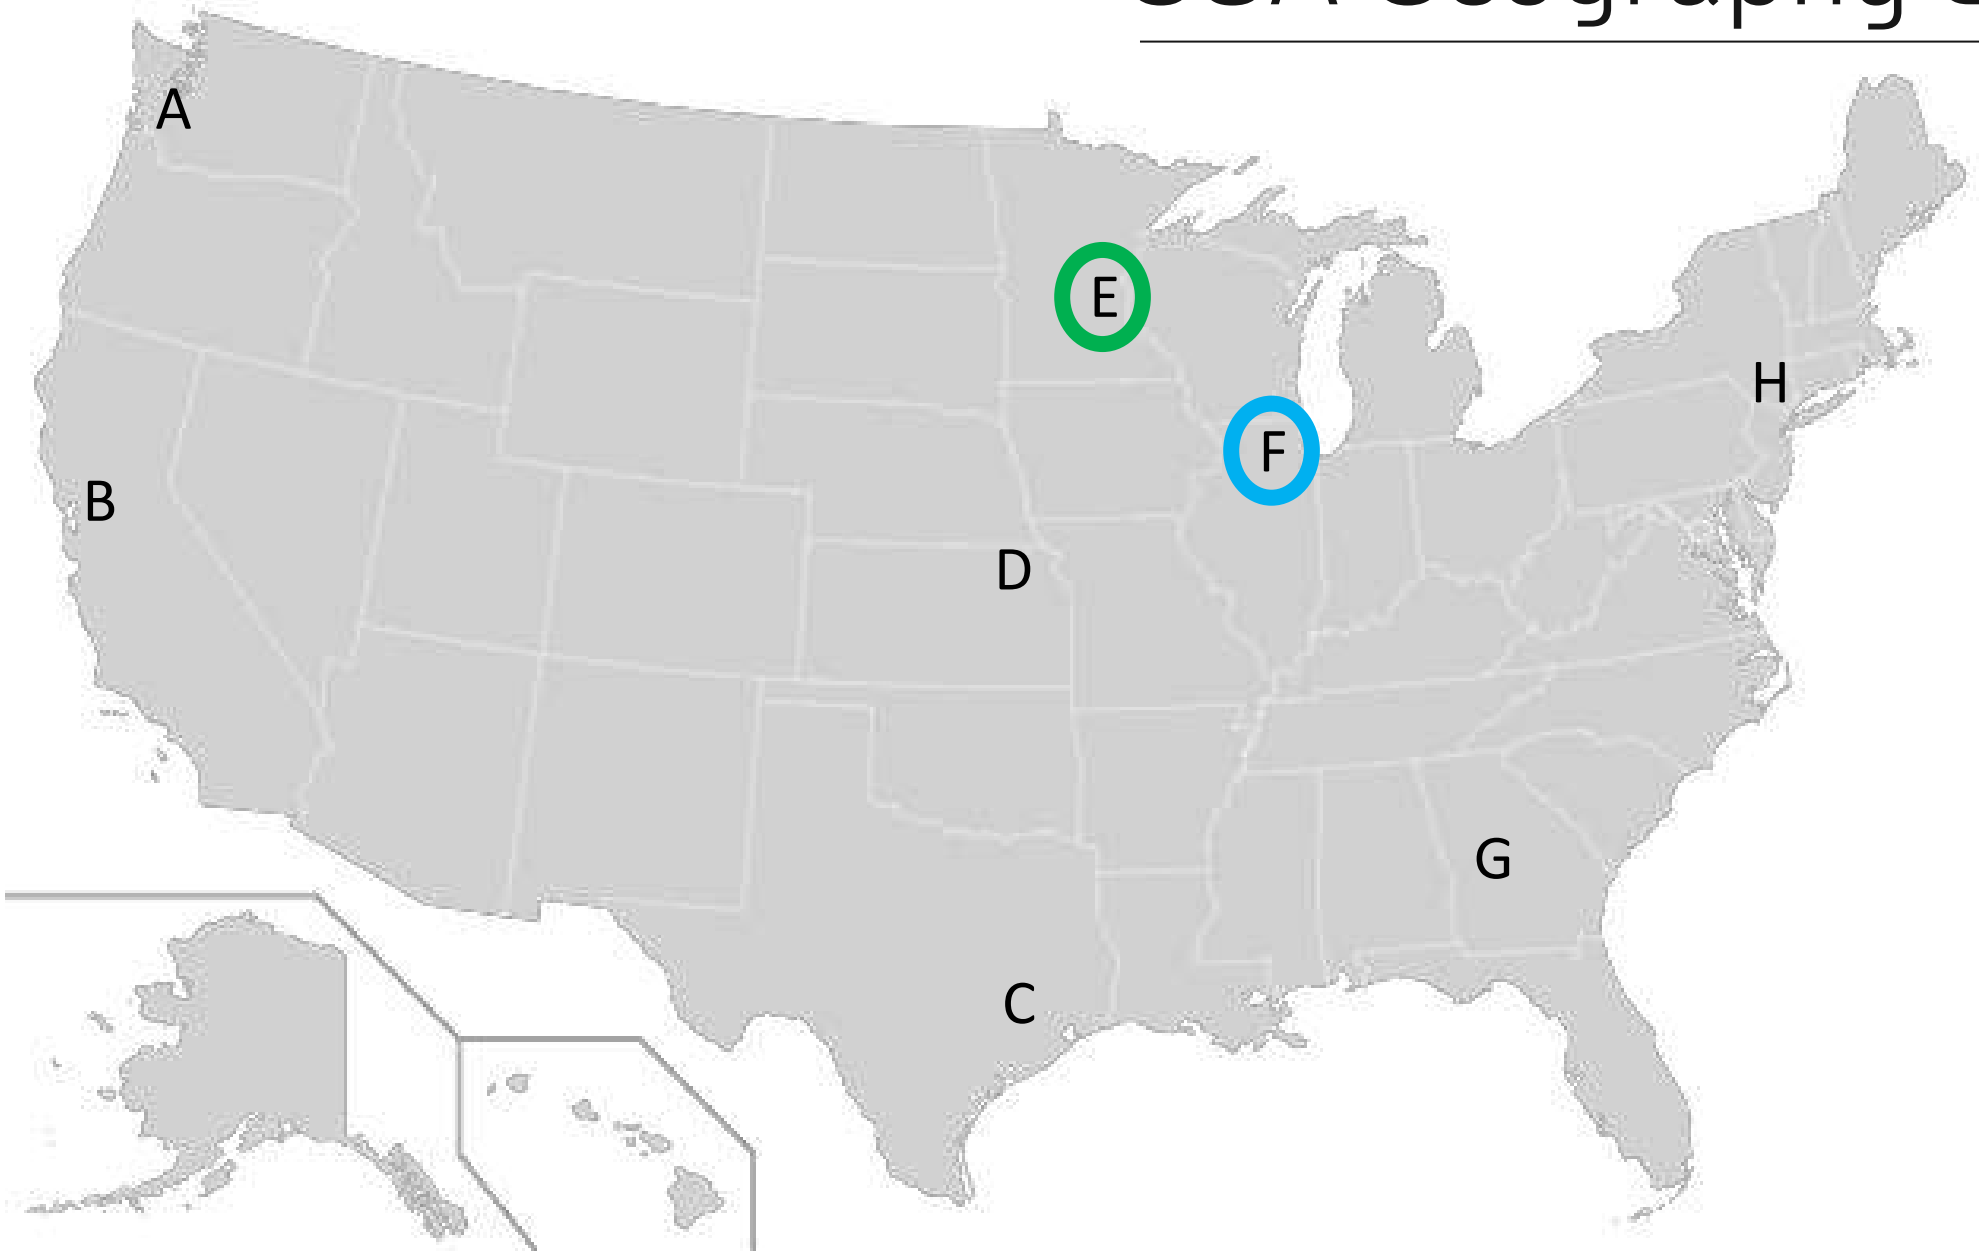

An Intro to Elias

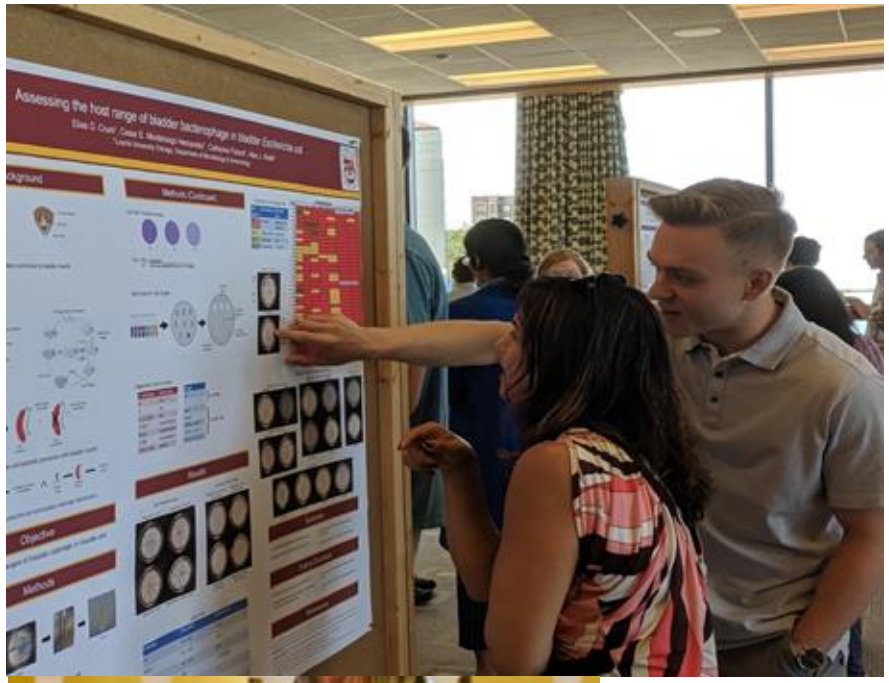

My Journey:

LOYOLA UNIVERSITY CHICAGO 1870 AD MAJOREM DEI GLORIAM

BSc + MSc - Bioinformatics

Prokaryotic Genomics

How can we increase **genomics** use in **medicine**?

PENGQUIN: PErsoNal Genome QUery IN healthcare and clinical practice

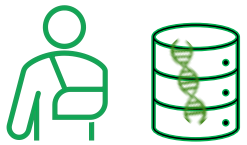

Genomic Data

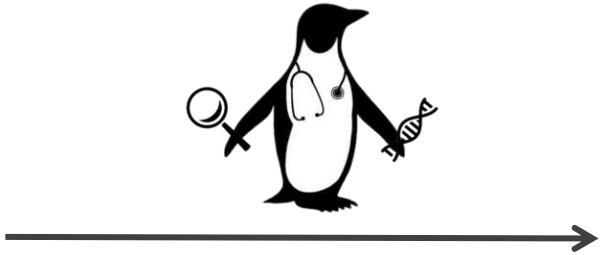

Clinical Care

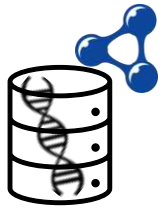

Semantic, Linked Data

Query Capabilities

Framework + App

Advisors: Dr. Gökhan Ertaylan[1], Dr. Bart Buelens[1], Dr. Ruben Taelman[2], Prof. Dr. Ruben Verborgh[2]

1 = vito 2 = GHENT UNIVERSITY

PENGQUIN Goals

Data Storage

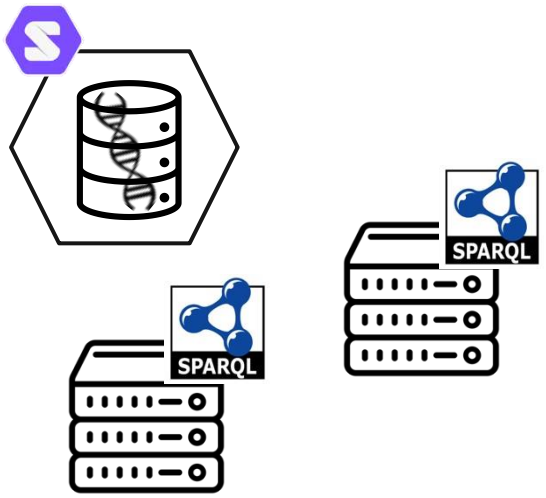

Genomic Data Representation

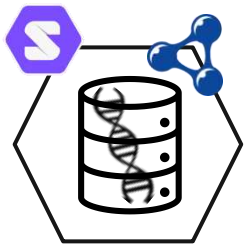

Question Answering

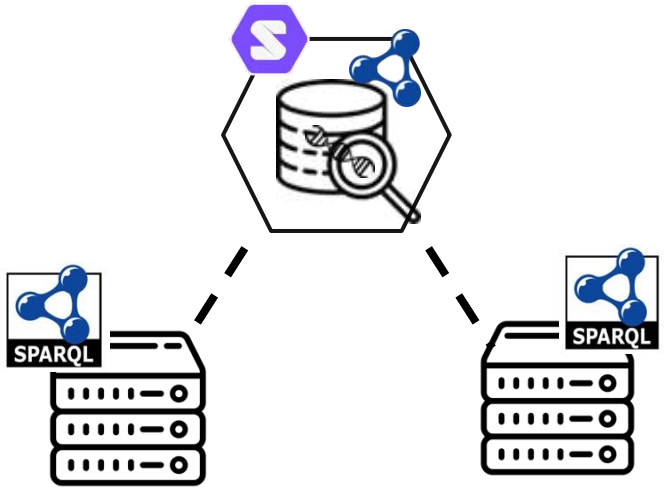

Framework
& web application

PENGQUIN Goals

Improve genomic data USAGE and SCALABILITY

Improve genomic data store accessibility (while maintaining privacy)

Represent genomic data semantically as Linked Data

Query genomic data and linkages

Beacon

Connect to existing initiatives / applications

Proof-of-concept

Genome Pods

Develop a framework and web application

Talk Outline

Decentralized Data Storage *(CHIST-ERA TRIPLE Project)*

- Introduction to Solid Pods
- Introduction to TRIPLE Project
- The Solid Cockpit web application (and maybe a use-case example if we have time)

Genomic Variant Data Semantification *(Semantic VCF)*

- Why Semantic VCF Data?
- Existing vocabularies/tools

Federated Querying *(Real World Federated Query Survey)*

- Motivation + Challenges
- Progress and preliminary findings
- Future aims

The First Act – Private Data Storage

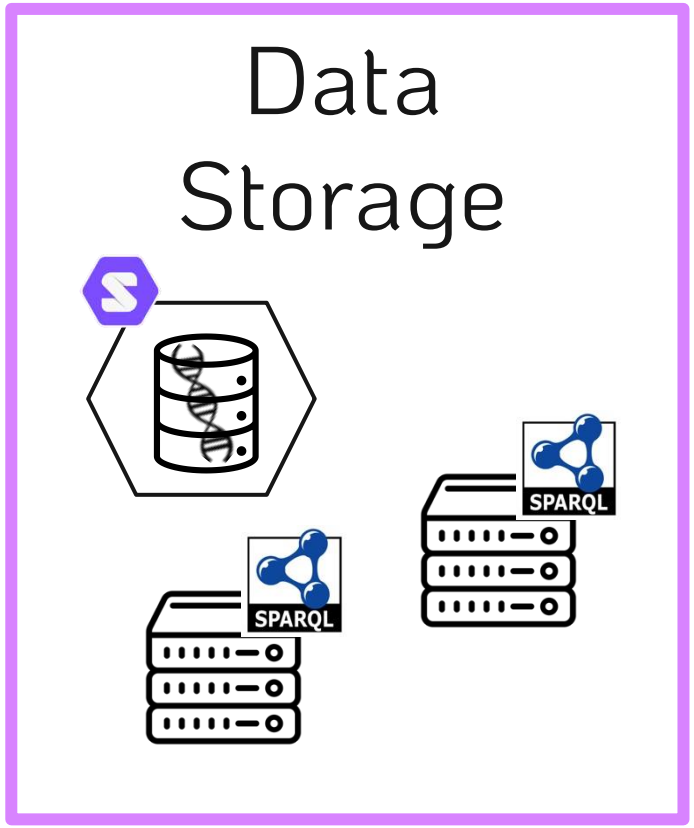

CHIST-ERA TRIPLE Project

Swiss Institute of Bioinformatics

GHENT UNIVERSITY

ÚOCHB AV ČR IOCB PRAGUE

chist-era

Address current **challenges** of using **RDF resources**
to bring the **utility** of these resources to a **broader group of researchers**

What is a Solid Pod?

Secure, Decentralized Data Storage Specification

- Can store files, RDF data, and Linked Data
- Provides control policies for stored data
- Pod exists on a Server (accessible to the Web)

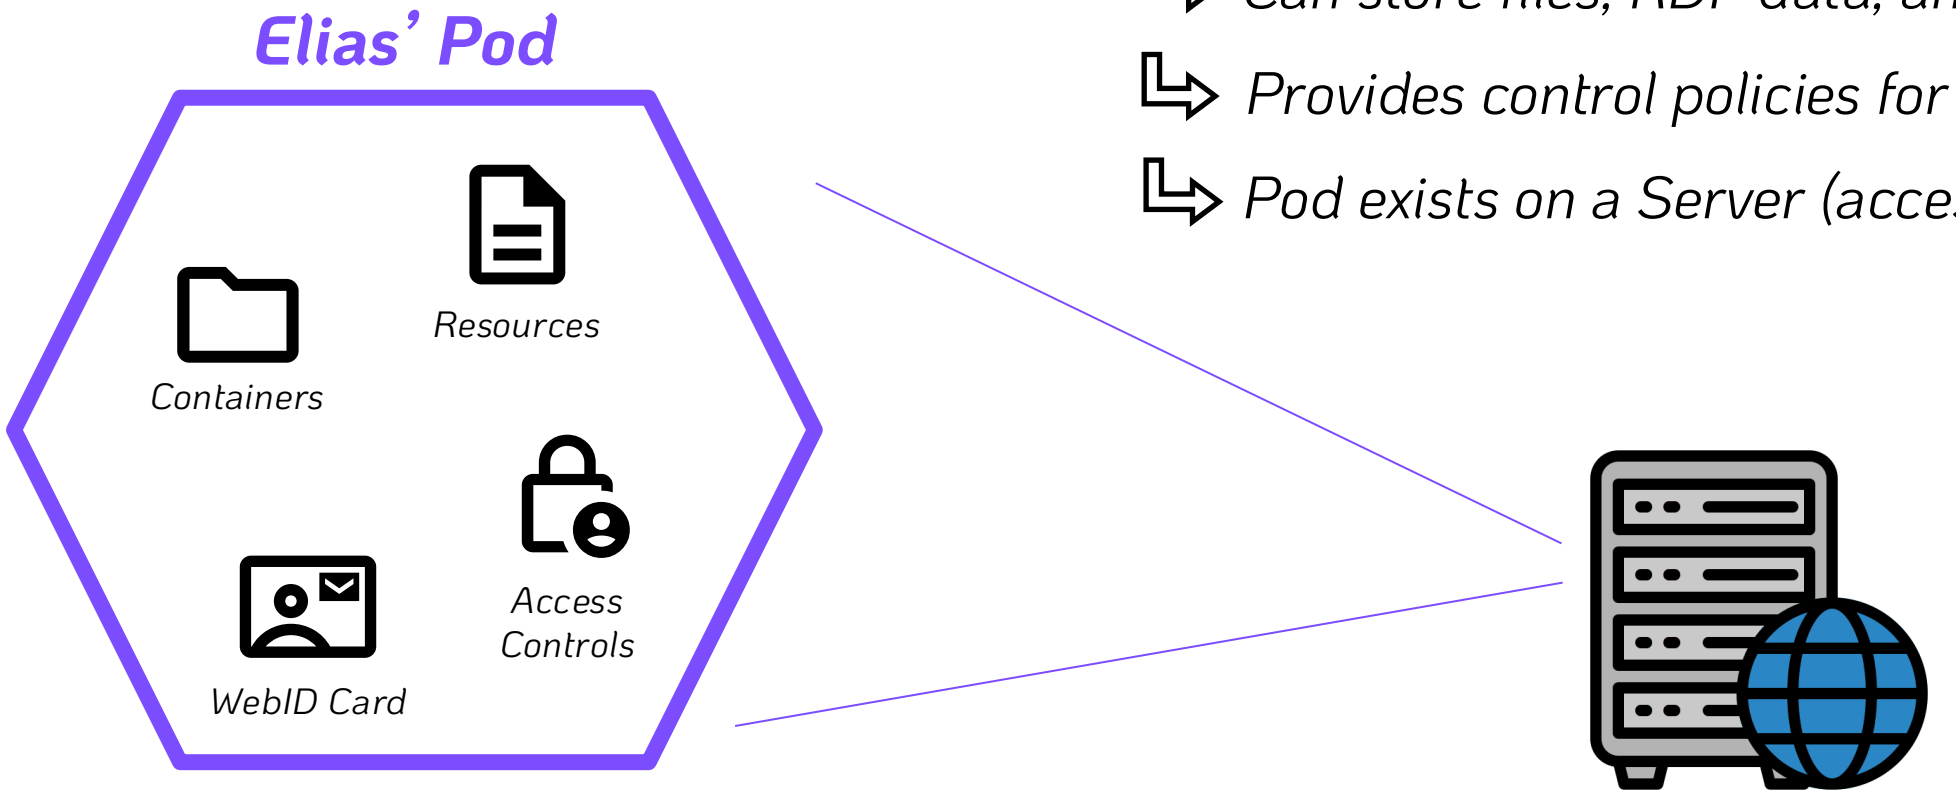

https://solidproject.org/about

Solid's Architecture

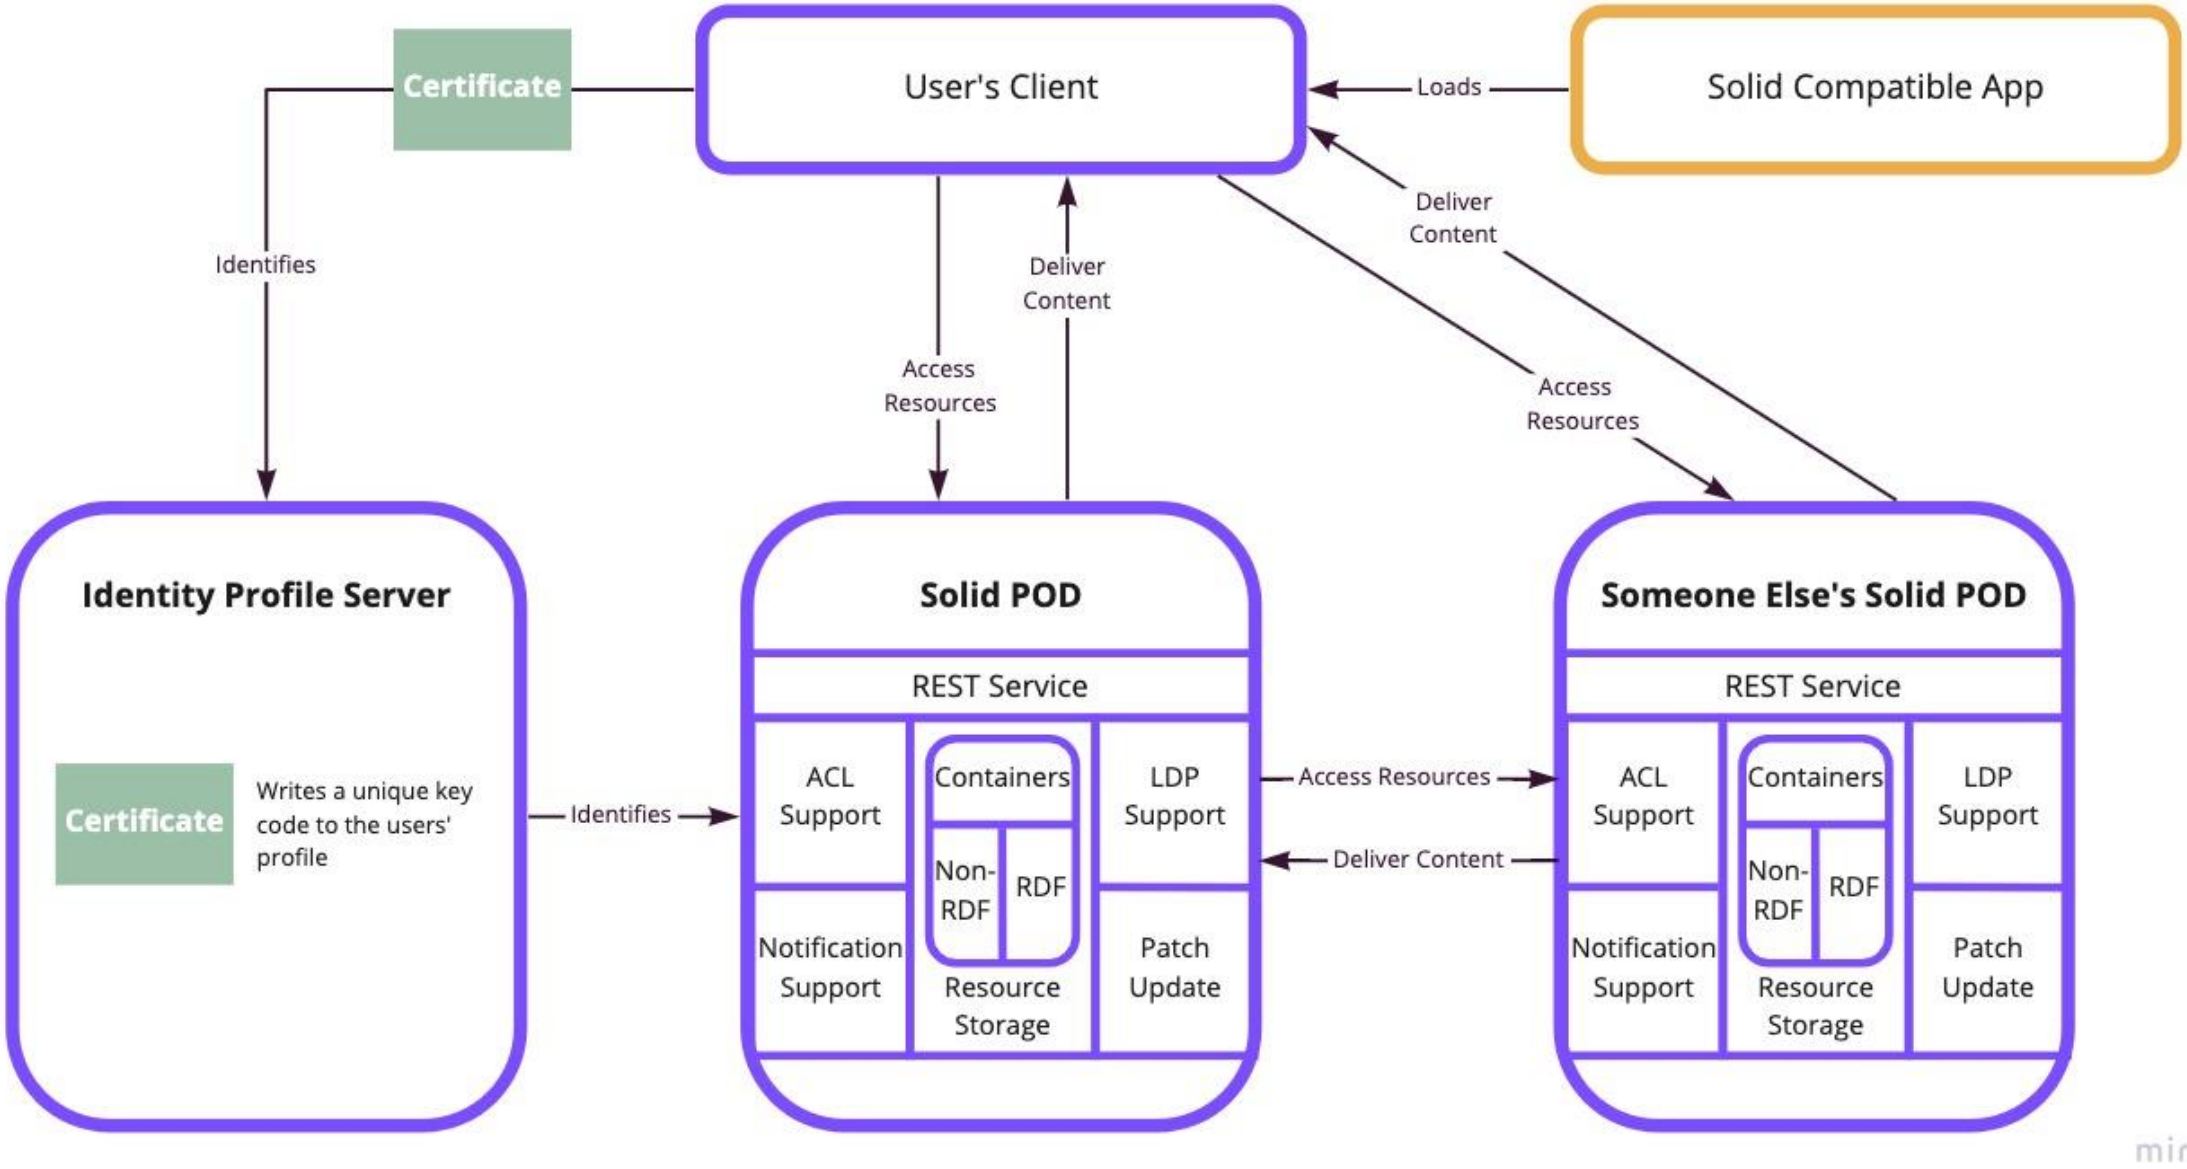

https://solid.georgetown.domains/technical-component/

Solid Pod vs SPARQL Endpoint?

Solid Pod

+ decentralized storage strategy
+ can store both file data and RDF data
+ granular privacy policy capabilities

- native querying capabilities limited
- not designed for large-scale data
- no computation attached to data store

SPARQL Endpoint

- centralized storage strategy
- triple store (only RDF)
- limited privacy capabilities

+ designed for high-performance querying
+ excels with large datasets
+ computation directly attached to data store

How to query a Solid Pod?

A Query Engine

- Query computation client-side
- Allows for SPARQL query execution "over" Solid Pod API
- Linked Data traversal + heterogeneous result joining capabilities

SPARQL Query

```
SELECT DISTINCT ?p {
    ?s ?p ?o
} LIMIT 10
```

+ Source:
Elias' Pod

Results Binding Stream
SPARQL+JSON

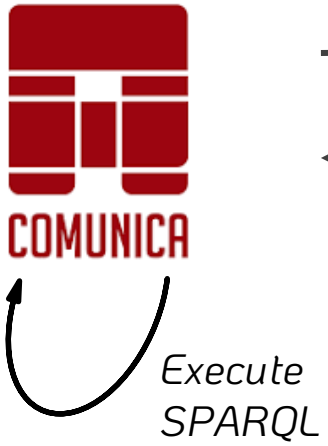

Request

Download

Elias' Pod

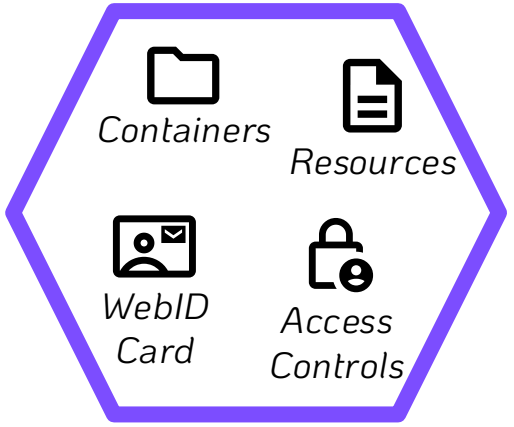

https://comunica.dev/about

Solid Pod Implementations?

BBC MyPDS App

Flemish Data Utility Company

health data vaults

Flemish Smart Dataspace

Solid Lab Use-cases

CHIST-ERA TRIPLE Project

Standardized documentation and resource statistics

My Contributions

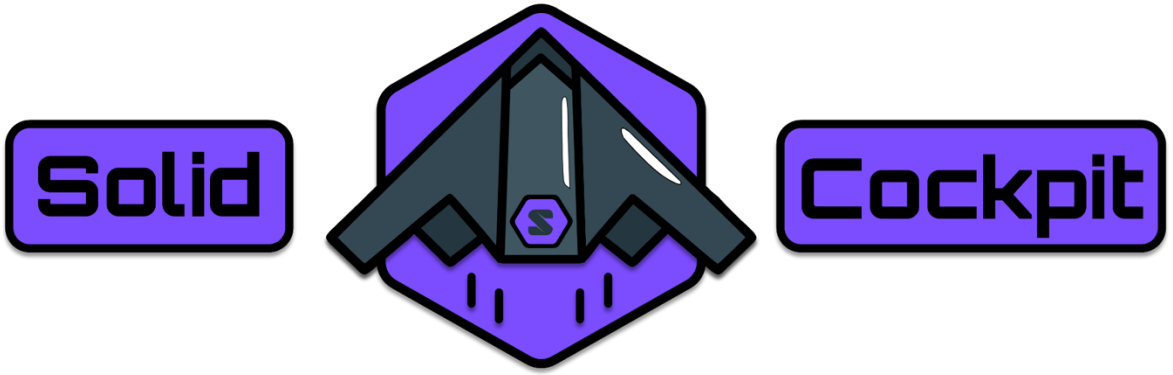

knowledgeonwebscale.
github.io/solid-cockpit/

A Web Application to:

- Upload Data to a Solid Pod
- Browse contents of a Solid Pod
- Edit Privacy of (and share) data in a Solid Pod
- Query Data in Solid Pods + SPARQL Endpoints
  - Implement a query cache for storing queries and results

The Second Act – Semantifying Genomic Variant Data

Data Storage

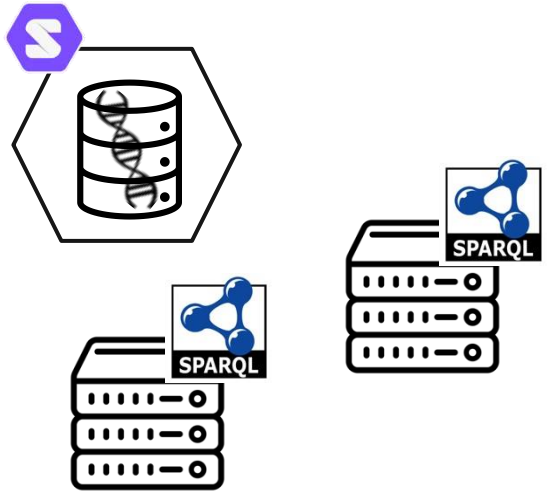

Genomic Data Representation

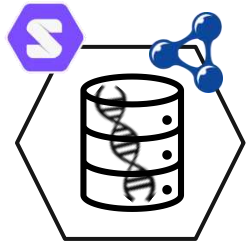

Representing Genomic Variant Data Semantically

Genomic Data, Health Care, and Semantics?

Pharmacogenomics Use-case (E.g. Warfarin):

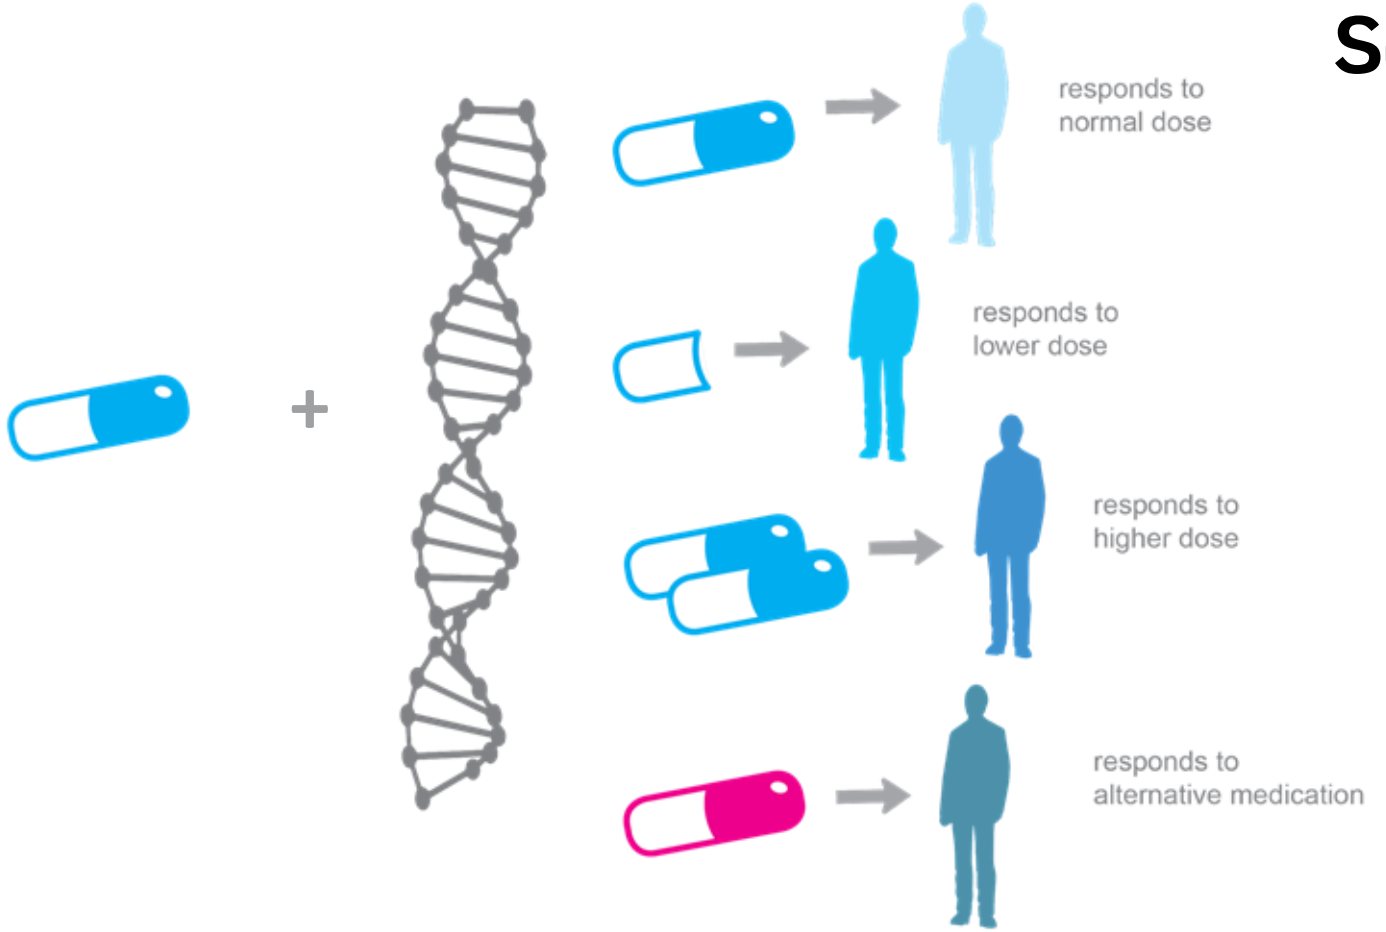

https://alphagenomix.com/empowering-personalized-medicine/

Semantic Genomic Variant Data?

Research:

↳ **F.A.I.R.** –
Findable. Accessible.
Interoperable. Reproducible.

Clinic:

↳ Data Linking
Data Provenance
Data Interoperability

Semantic VCF – *Sequencing Data Formats Intro*

FASTQ ✚ **FASTA** → **SAM/BAM** → **VCF**

Raw Sequencing Outputs

Contiguous DNA sequence

FASTQ reads aligned to a FASTA reference

Areas in SAM where sample and reference differ

```
1  >gi|568336023|gb|CM000663.2| Homo sapiens chromosome 1, GRCh
2  CCTAACCCTAACCCTAACCCTAACCCTAACCCTAACCCTAACCCTAACCCTAACCCTAAC
3  ACCCTAACCCTAACCCTAACCCTAACCCAACCCTAACCCTAACCCTAACCCTAACCCTAA
4  TAACCCTAACCCTAACCCTAACCCTAACCTAACCCTAACCCTAACCCTAACCCTAACCCT
5  CTAACCCTAACCCCTAACCCTAACCCTAAACCCTAAACCCTA
6  CAACCCCAACCCCAACCCCAACCCCAACCCCAACCCTAACCC
7  CCTAACCCTAACCCTAACCCTAACCCTAACCCCTAACCCCTA
8  TAACCCTAACCCCTAACCCTAACCCTAACCCTAACCCTCGCG
9  TCTGACCTGAGGAGAACTGTGCTCCGCCTTCAGAGTACCACC
10 CGCCCTCGCGGTGCTCTCCGGGTCTGTGCTGAGGAGAACGCA
11 CGGCGCAGGCGCAGAGAGGCGCGCCGCGCCGGCGCAGGCGCA
12 AGAGAGGCGCGCCGCGCCGGCGCAGGCGCAGAGAGGCGCGCC
13 CGCGCCGGCGCAGGCGCAGACACATGCTAGCGCGTCGGGGTG
14 GCCGCGCCGGCGCAGGCGCAGAGACACATGCTACCGCGTCCA
15 AGGCGCACCGCGCCGGCGCAGGCGCAGAGACACATGCTAGCG
16 GCAGAGACGCAAGCCTACGGGCGGGGGTTGGGGGGGCGTGTG
17 GCTGGGGCGGGGGGAGGGTGGCGCCGTGCACGCGCAGAAACT
18 GGTAGAACCTCAGTAATCCGAAAAGCCGGGATCGACCGCCCC
19 GCTTGCTCACGGTGCTGTGCCAGGGCGCCCCCTGCTGGCGAC
20 GAGTGGTGGCCAGCGCCCCCTGCTGGCGCCGGGGCACTGCAG
```

```
@FORJUSP02
CCGTCAATTC
+
AAAAAAAAAA
```

Q scores (as ASCII ch

Base=T, Q=':'=25

```
@HD     VN:1.0  SO
@SQ     SN:chr20
@PG     ID:TopHat
lign-edit-dist 2 -
20 /data/user446/m
HWI-ST1145:74:C101
       CCGTGTTTAAA
C    BBDCCDDCCDDDD
    AS:i:-15
HWI-ST1145:74:C101
       TGCTGGATCAT
G    DCDDDDEDDDDDD
    AS:i:-16
HWI-ST1145:74:C101DACXX:7:1204:14760:4030       16       chr20    270877  50       100M     *       0       0
       GGCTTTATTGGTAAAAAAGGAATAGCAGATTTAATCAGAAATTCCCACCTGGCCCAGCAGCACCAACCAGAAAGAAGGGAAGAAGACAGGAAAAAACCA
C    DDDDDDDDDCCDDDDDDDDDDEEEEEEEFFFEFFEGHHHHFGDJJIHJJIJIJJJIIIIGGFJJIHIIIIJJJJJJIGHHFAHGFHJHFGGHFFFDD@BB
    AS:i:-11        XM:i:2  XO:i:0  XG:i:0  MD:Z:0A85G13    NM:i:2  NH:i:1
HWI-ST1145:74:C101DACXX:7:1210:11167:8699       0       chr20    271218  50       50M4700N50M     *       0
       0       GTGGCTCTTCCACAGGAATGTTGAGGATGACATCCATGTCTGGGGTGCACTTGGGTCTCCGAAGCAGAACATCCTCAAATATGACCTCTCG
accepted_hits.sam
```

```
##fileformat=VCFv4.3
##fileDate=20090805
##source=myImputationProgramV3.1
##reference=file:///seq/references/1000GenomesPilot-NCBI36.fasta
##contig=<ID=20,length=62435964,assembly=B36,md5=f126cdf8a6e0c7f379d618ff66beb2da,species="Homo sapiens",taxonomy=x>
##phasing=partial
##INFO=<ID=NS,Number=1,Type=Integer,Description="Number of Samples With Data">
##INFO=<ID=DP,Number=1,Type=Integer,Description="Total Depth">
##INFO=<ID=AF,Number=A,Type=Float,Description="Allele Frequency">
##INFO=<ID=AA,Number=1,Type=String,Description="Ancestral Allele">
##INFO=<ID=DB,Number=0,Type=Flag,Description="dbSNP membership, build 129">
##INFO=<ID=H2,Number=0,Type=Flag,Description="HapMap2 membership">
##FILTER=<ID=q10,Description="Quality below 10">
##FILTER=<ID=s50,Description="Less than 50% of samples have data">
##FORMAT=<ID=GT,Number=1,Type=String,Description="Genotype">
##FORMAT=<ID=GQ,Number=1,Type=Integer,Description="Genotype Quality">
##FORMAT=<ID=DP,Number=1,Type=Integer,Description="Read Depth">
##FORMAT=<ID=HQ,Number=2,Type=Integer,Description="Haplotype Quality">
#CHROM POS     ID        REF    ALT     QUAL FILTER INFO                              FORMAT      NA00001        NA00002        NA00003
20     14370   rs6054257 G      A       29   PASS   NS=3;DP=14;AF=0.5;DB;H2           GT:GQ:DP:HQ 0|0:48:1:51,51 1|0:48:8:51,51 1/1:43:5:.,.
20     17330   .         T      A       3    q10    NS=3;DP=11;AF=0.017               GT:GQ:DP:HQ 0|0:49:3:58,50 0|1:3:5:65,3   0/0:41:3
20     1110696 rs6040355 A      G,T     67   PASS   NS=2;DP=10;AF=0.333,0.667;AA=T;DB GT:GQ:DP:HQ 1|2:21:6:23,27 2|1:2:0:18,2   2/2:35:4
20     1230237 .         T      .       47   PASS   NS=3;DP=13;AA=T                   GT:GQ:DP:HQ 0|0:54:7:56,60 0|0:48:4:51,51 0/0:61:2
20     1234567 microsat1 GTC    G,GTCT  50   PASS   NS=3;DP=9;AA=G                    GT:GQ:DP    0/1:35:4       0/2:17:2       1/1:40:3
```

What does a VCF File Look like?

Example

VCF header

```
##fileformat=VCFv4.0
##fileDate=20100707
##source=VCFtools
##reference=NCBI36
##INFO=<ID=AA,Number=1,Type=String,Description="Ancestral Allele">
##INFO=<ID=H2,Number=0,Type=Flag,Description="HapMap2 membership">
##FORMAT=<ID=GT,Number=1,Type=String,Description="Genotype">
##FORMAT=<ID=GQ,Number=1,Type=Integer,Description="Genotype Quality (phred score)">
##FORMAT=<ID=GL,Number=3,Type=Float,Description="Likelihoods for RR,RA,AA genotypes (R=ref,A=alt)">
##FORMAT=<ID=DP,Number=1,Type=Integer,Description="Read Depth">
##ALT=<ID=DEL,Description="Deletion">
##INFO=<ID=SVTYPE,Number=1,Type=String,Description="Type of structural variant">
##INFO=<ID=END,Number=1,Type=Integer,Description="End position of the variant">
```

Body

#CHROM	POS	ID	REF	ALT	QUAL	FILTER	INFO	FORMAT	SAMPLE1	SAMPLE2
1	1	.	ACG	A,AT	.	PASS	.	GT:DP	1/2:13	0/0:29
1	2	rs1	C	T,CT	.	PASS	H2;AA=T	GT:GQ	0|1:100	2/2:70
1	5	.	A	G	.	PASS	.	GT:GQ	1|0:77	1/1:95
1	100	.	T	<DEL>	.	PASS	SVTYPE=DEL;END=300	GT:GQ:DP	1/1:12:3	0/0:20

Mandatory header lines

Optional header lines (meta-data about the annotations in the VCF body)

Reference alleles (GT=0)

Alternate alleles (GT>0 is an index to the ALT column)

Deletion

SNP

Large SV

Insertion

Other event

Phased data (G and C above are on the same chromosome)

https://davetang.github.io/learning_vcf_file/

VCF Spec: https://samtools.github.io/hts-specs/VCFv4.2.pdf

Representing VCF Semantically?

GA4GH Variant Schema:

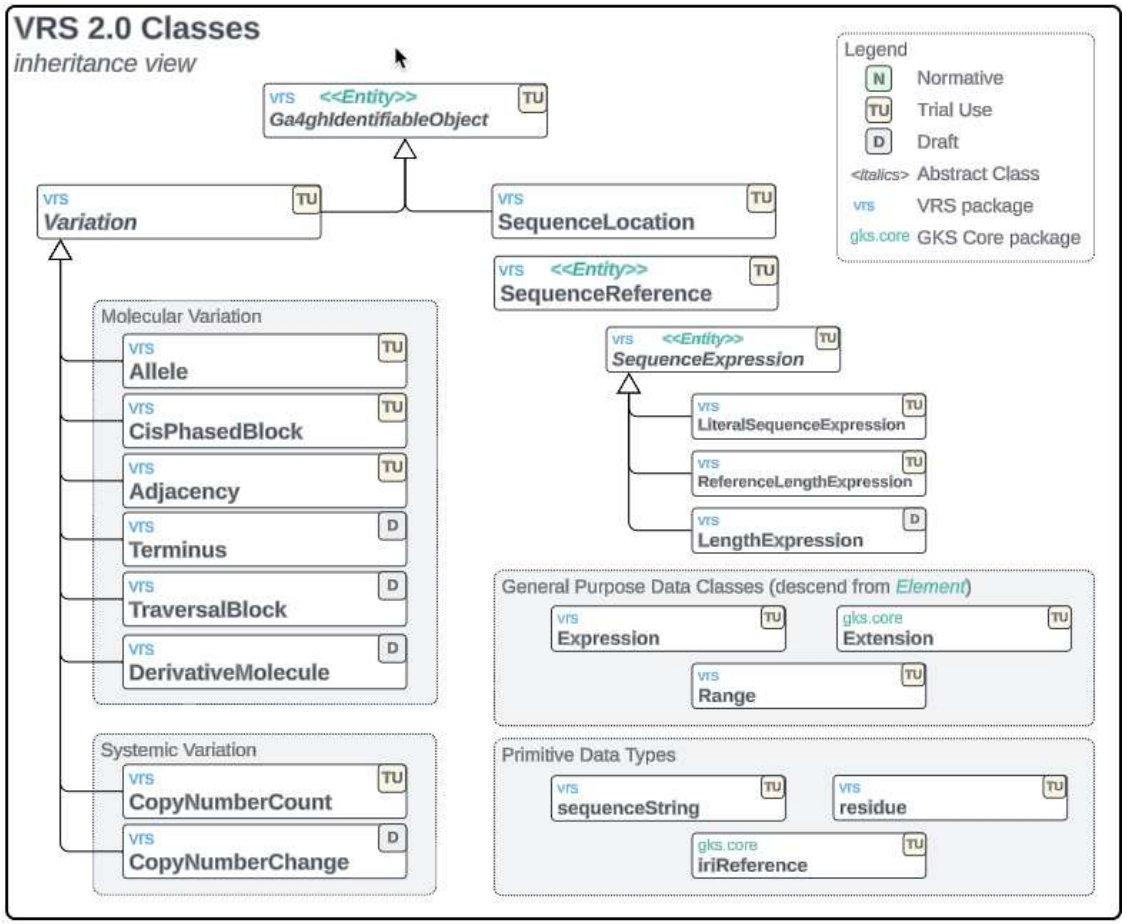

HERO Genomics Ontology:

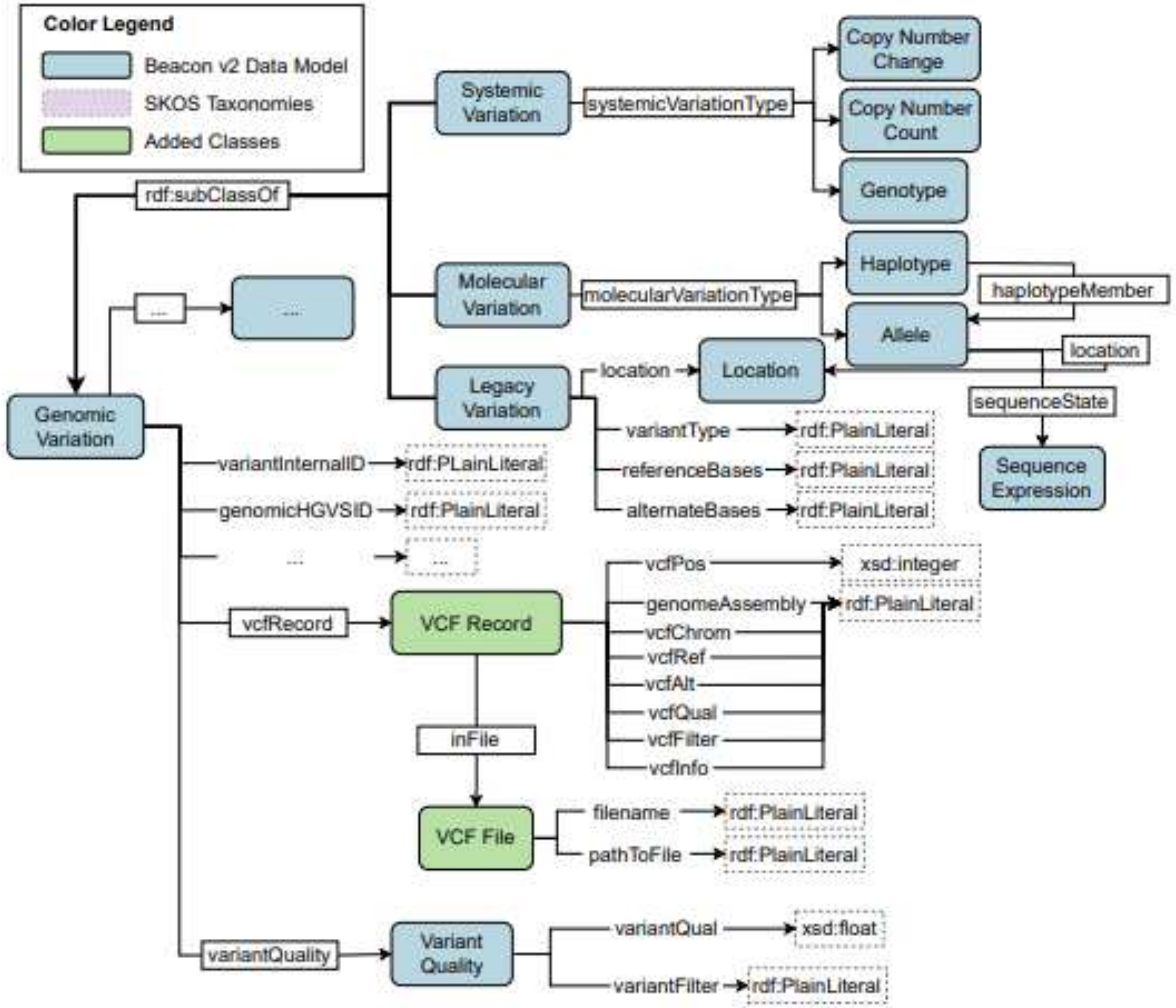

https://vrs.ga4gh.org/en/2.0.0-ballot.2024-11/appendices/class_diagram.html#classdiagram

https://www.dei.unipd.it/~silvello/papers/2025_MCS_SWAT4HCLS.pdf

Semantic VCF - Future Aims

Construct VCF Knowledge Graph

rml.io

OR

RocketRML

Experiment with SPARQL Querying and Data Linking

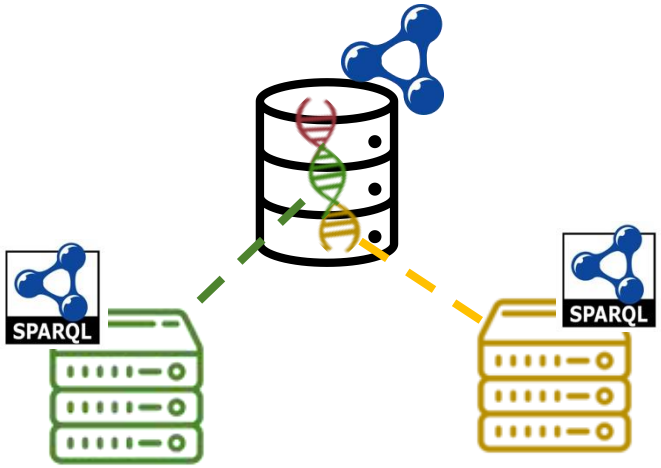

The Third Act – SPARQL Federated Querying

Data Storage

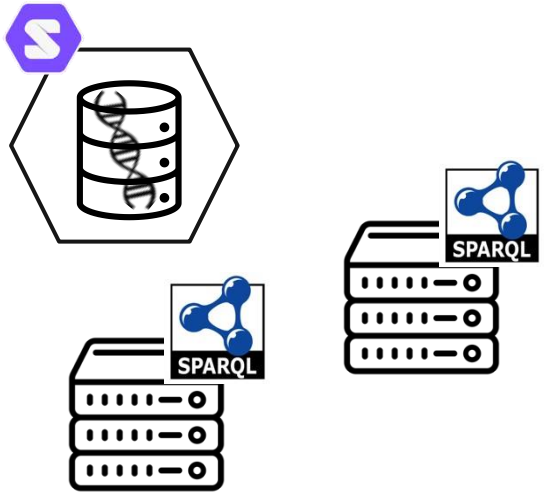

Genomic Data Representation

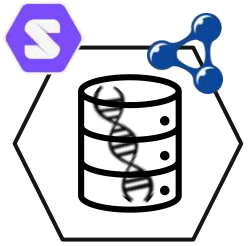

Question Answering

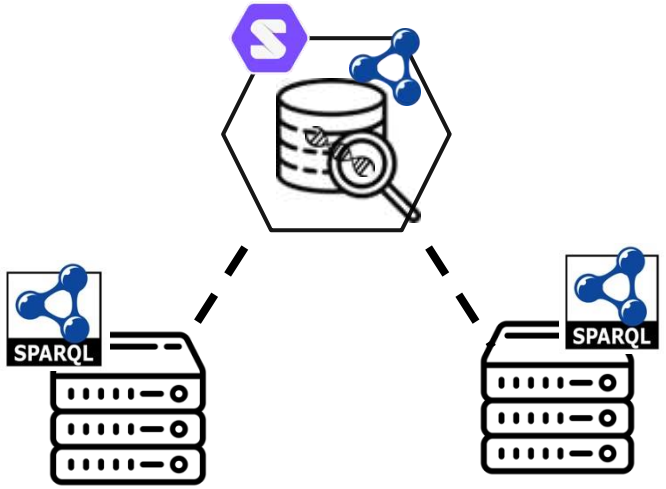

Real World Federated Query Survey

Swiss Institute of Bioinformatics

GHENT UNIVERSITY

Motivation:

Federated querying for answering questions over **multiple heterogeneous data sources**

- ? Performance is currently ~*variable*~
  - ↳ We want to detail *WHY* performance varies so much && suggestions for improvement

- ? Assess federation algorithms on <u>REAL endpoints</u>
  - ↳ ***<u>FedX</u>*** – uses `ASK` queries to inform source-selection
  - ↳ ***<u>SPLENDID</u>*** – `VoID` descriptions to inform source-selection

Real World Federated Query Survey - *Design*

SPARQL Endpoints:

"Real World" Federated Queries: https://sib-swiss.github.io/sparql-examples/

SPARQL Query Engine: https://comunica.dev/

Example Federated Query

Query: 027-biosodafrontend

Question: (*Bgee* | *UniProt*)

Retrieve the *genes* expressed in the *human pancreas*

and

those *genes*' *annotations*

https://www.bgee.org/sparql/

```
PREFIX rdfs: <http://www.w3.org/2000/01/rdf-schema#>
PREFIX up: <http://purl.uniprot.org/core/>
PREFIX obo: <http://purl.obolibrary.org/obo/>
PREFIX genex: <http://purl.org/genex#>
PREFIX lscr: <http://purl.org/lscr#>
PREFIX orth: <http://purl.org/net/orth#>
SELECT DISTINCT ?geneEns ?uniprot ?annotation_text {
    {
        SELECT ?geneEns {
            ?geneB a orth:Gene .
            ?geneB genex:isExpressedIn ?cond .
            ?cond genex:hasAnatomicalEntity ?anat .
            ?geneB lscr:xrefEnsemblGene ?geneEns .
            ?anat rdfs:label 'pancreas' .
            ?geneB orth:organism ?o .
            ?o obo:RO_0002162 ?taxon2 .
            ?taxon2 up:commonName 'human' .
        } LIMIT 100
    }
    SERVICE <https://sparql.uniprot.org/sparql> {
        ?uniprot rdfs:seeAlso / up:transcribedFrom ?geneEns .
        ?uniprot up:annotation ?annotation .
        ?annotation rdfs:comment ?annotation_text .
    }
}
```

visualization from: https://sib-swiss.github.io/sparql-editor/

Example obtained from: https://biosoda.expasy.org/build_biosodafrontend/

Real World Federated Query Survey – *[Preliminary] Findings 1*

Preliminary Results (endpoints and queries):

Federated querying is not "easy"

- Errors related to SPARQL endpoint availability

  Invalid SPARQL endpoint response from https://semantic.eea.europa.eu/sparql (HTTP status 404)

- Errors caused by SPARQL endpoint-related errors

  Invalid SPARQL endpoint response from https://sparql.orthodb.org/sparql (HTTP status 500)

- Errors related to SPARQL endpoint rate-limiting

  Invalid SPARQL endpoint response from https://query.wikidata.org/sparql (HTTP status 429): Rate limit exceeded

- Errors related to SPARQL endpoint time-outs

  Unexpected ""!"" at position 0 in state STOP

Real World Federated Query Survey – *[Preliminary] Findings 2*

Preliminary Results (algorithms):

Federated querying is not ”easy”

↳ Errors within Comunica

Actor urn:comunica:default:query-operation/actors#source requires an operation with source annotation.

fetch failed

Could not retrieve https://idsm.elixir-czech.cz/sparql/endpoint/chebi (HTTP status 400)

FATAL ERROR: Ineffective mark-compacts near heap limit Allocation failed - JavaScript heap out of memory

↳ Errors related to algorithms

Actor urn:comunica:default:/actors#rdf-join cannot complete join with Asym-hash actor.

Real World Federated Query Survey – *Future Work*

Future Aims:

Execute 1 month worth of queries

↳ 4 experimental runs per week

Vary time + part of week experiment is performed

Publication goals:

- In-Use Paper
- Journal Paper Extension (benchmark creation)

Algorithmic variation

- `VoID` vs `ASK` queries
- Various join algorithm combinations
- Rate-limiting
- `SERVICE` removal

Metrics Assessed

- Total execution time
- Query planning time
- Result set consistency
- # SPARQL requests
- # HTTP requests
- Individual query arrival times

For more detail see: https://docs.google.com/document/d/1U-RtYsJw9raw4U5Qe6BHfSh2VYGnUIkbT9-aSyJionU/edit?usp=sharing

The Finale – Unified System for Storage & Representation & Querying

Data Storage

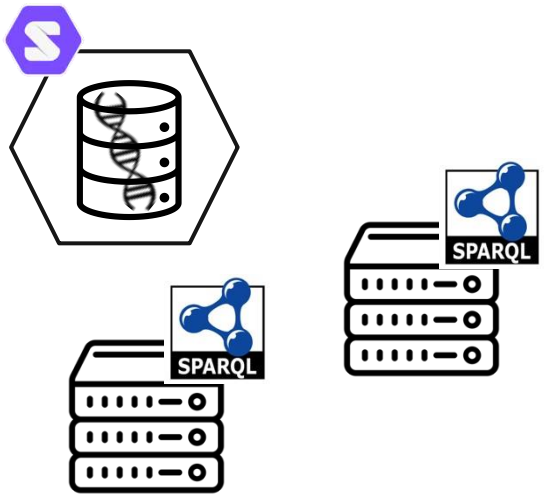

Genomic Data Representation

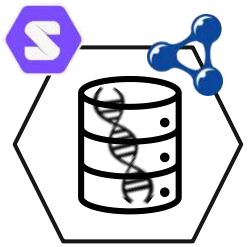

Question Answering

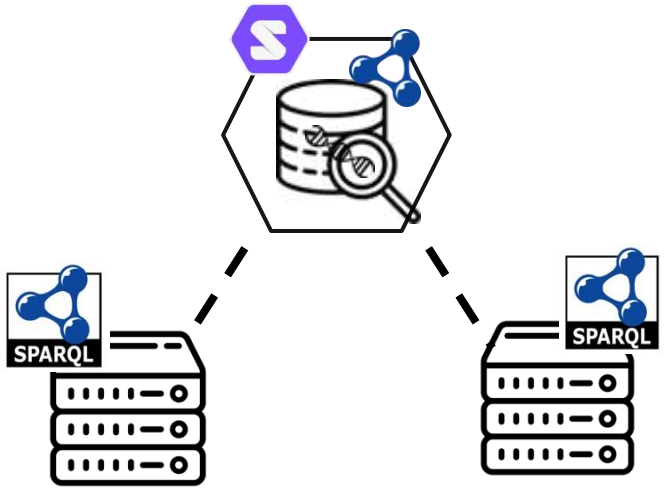

Framework
& web application

Thank you!

More PhD specifics

LinkedIn

GitHub

Supplemental Slides

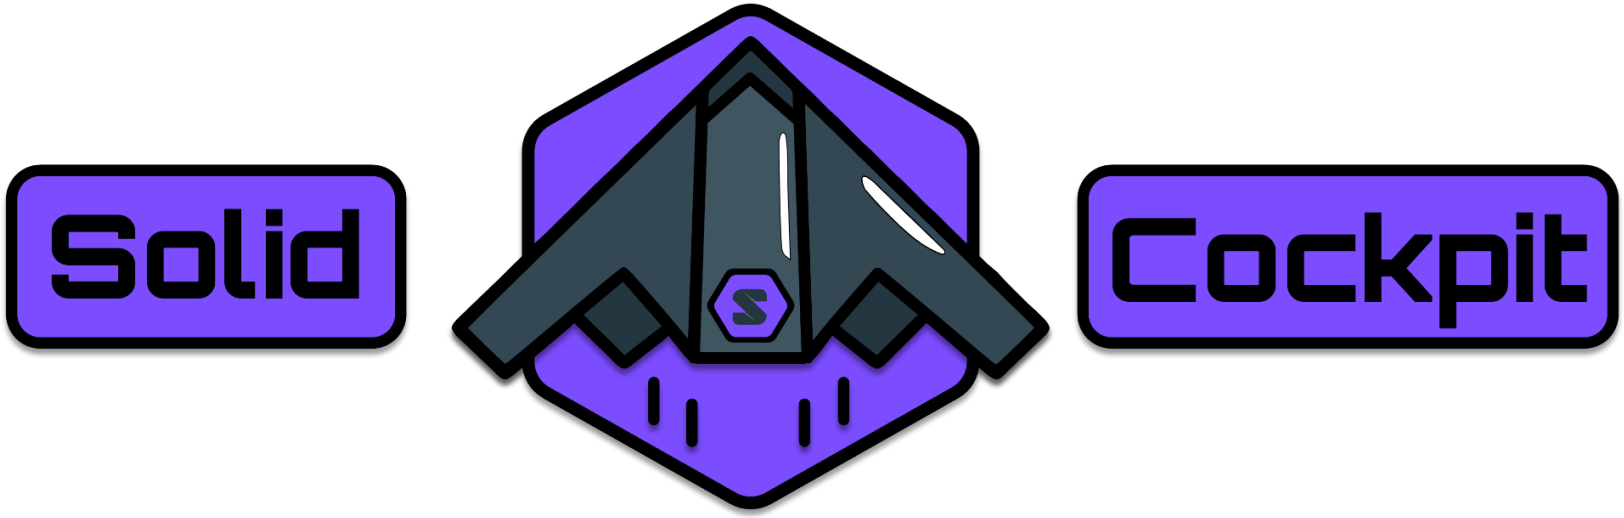

knowledgeonwebscale.github.io/solid-cockpit/

Genome Sequencing in Health Care

Sequence Generation Costs

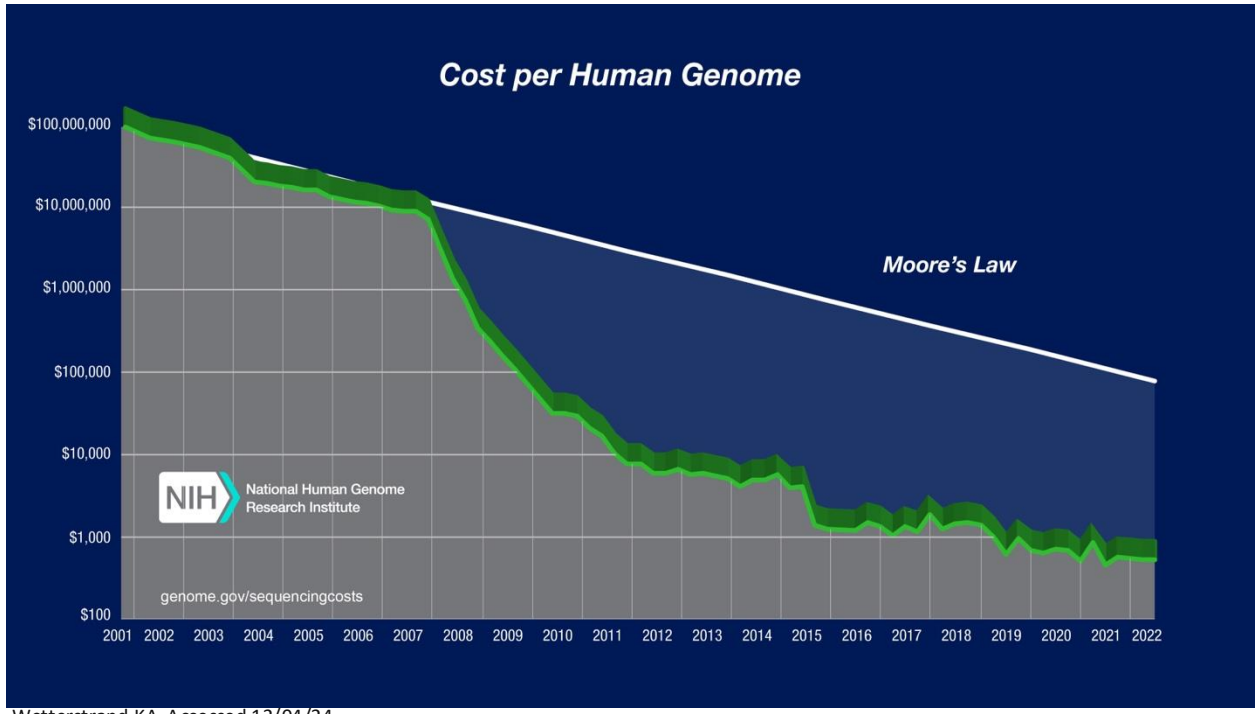

Wetterstrand KA. Accessed 12/04/24.

Currently getting close to $100 per genome[4]

Per-base cost of sequencing has been dropping by half every ~5 months

Per-base cost of sequence data storage is dropping by half every ~14 months

Use-cases

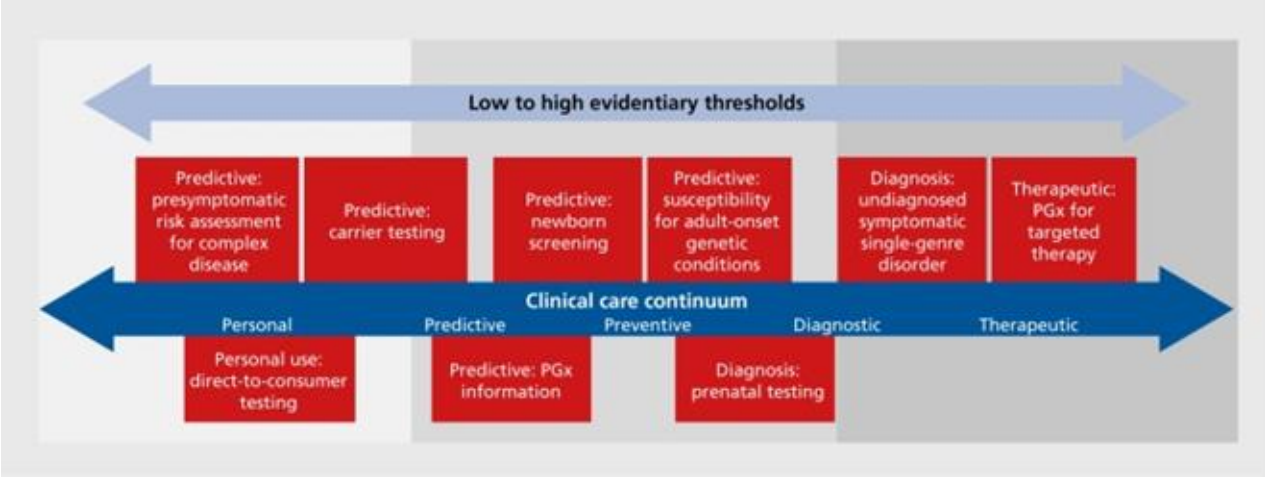

10.31887/DCNS.2016.18.3/jkrier

Examples include:

- Medication prescription (pharmacogenomics)
- Genetic disease screening
- Cancer diagnosis & treatment

Variant Call Format (VCF)

- tab delimited .txt file

- Rows represent individual mutations in a sequence (when compared to reference)

- Columns store information about that mutation

#CHROM	POS	ID	REF	ALT	QUAL	FILTER	INFO	FORMAT

```
##fileformat=VCFv4.1
##fileDate=20140930
##source=23andme2vcf.pl https://github.com/arrogantrobot/23
##reference=file://23andme_v3_hg19_ref.txt.gz
##FORMAT=<ID=GT,Number=1,Type=String,Description="Genotype"
#CHROM  POS ID  REF ALT QUAL    FILTER  INFO    FORMAT  GEN
chr1    82154   rs4477212   a   .   .   .   .   GT  0/0
chr1    752566  rs3094315   g   A   .   .   .   GT  1/1
chr1    752721  rs3131972   A   G   .   .   .   GT  1/1
chr1    798959  rs11240777  g   .   .   .   .   GT  0/0
chr1    800007  rs6681049   T   C   .   .   .   GT  1/1
chr1    838555  rs4970383   c   .   .   .   .   GT  0/0
chr1    846808  rs4475691   C   .   .   .   .   GT  0/0
chr1    854250  rs7537756   A   .   .   .   .   GT  0/0
chr1    861808  rs13302982  A   G   .   .   .   GT  1/1
chr1    873558  rs1110052   G   T   .   .   .   GT  1/1
chr1    882033  rs2272756   G   A   .   .   .   GT  0/1
chr1    888659  rs3748597   T   C   .   .   .   GT  1/1
chr1    891945  rs13303106  A   G   .   .   .   GT  0/1
```

Personal Genome Sequence Data

Human genome data is big…

~3.2 BILLION base pairs ----> ~1 GB

Sequencing output file ----> ~200 GB

Variant Call Format (***VCF***) file ----> ~200 MB

- Individual humans exhibit ~3 million mutations
- VCF file **records only those differences**
- Flat file

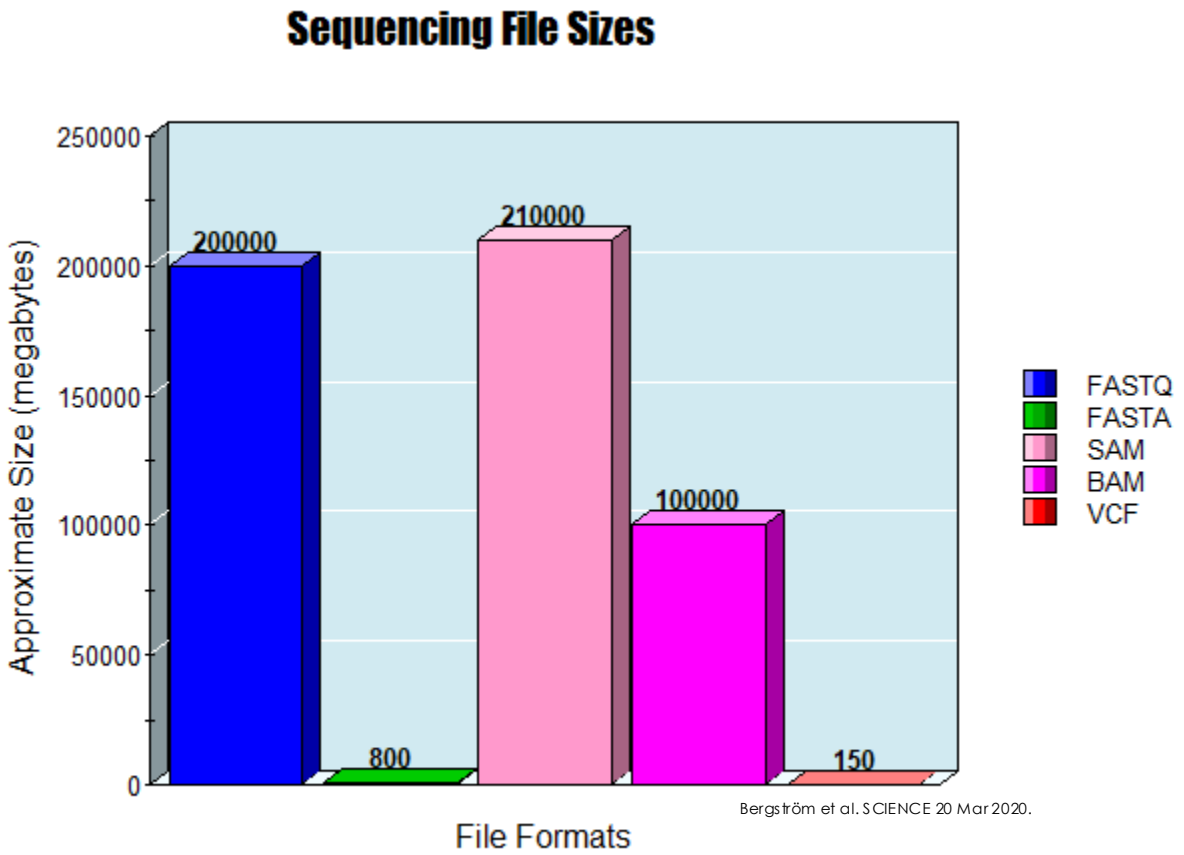

Bergström et al. SCIENCE 20 Mar 2020.

VCF file semantic representation

- **~ >27 MILLION** Triples when represented as RDF
- Ontology for conversion
- Header Dictionary Triples (for compression)

Scaling Genome Data Usage

Research genomic data sharing initiatives:

FHIR

Beacon

OHDSI

OBSERVATIONAL HEALTH DATA SCIENCES AND INFORMATICS

State-of-the-art for **clinical** genomic data sharing?

GDPR (EU) / HIPPA (US)
EU Health Data Spaces → Not much...

Experimental VCF Data

Primary dataset from publicly available repositories (>100 genomes)

- Personal Genome Project (https://my.pgp-hms.org/public_genetic_data)
- NIST's Genome in a Bottle public dataset (https://data.nist.gov/od/id/mds2-2336)
- Illumina Platinum Genomes repository (https://emea.illumina.com/platinumgenomes.html)

If more are needed (>1000 [research grade] genomes):

- 1000 genomes project (https://www.internationalgenome.org/)
- NCBI's ClinVar (https://www.ncbi.nlm.nih.gov/clinvar/)

Semantic Genome Data?

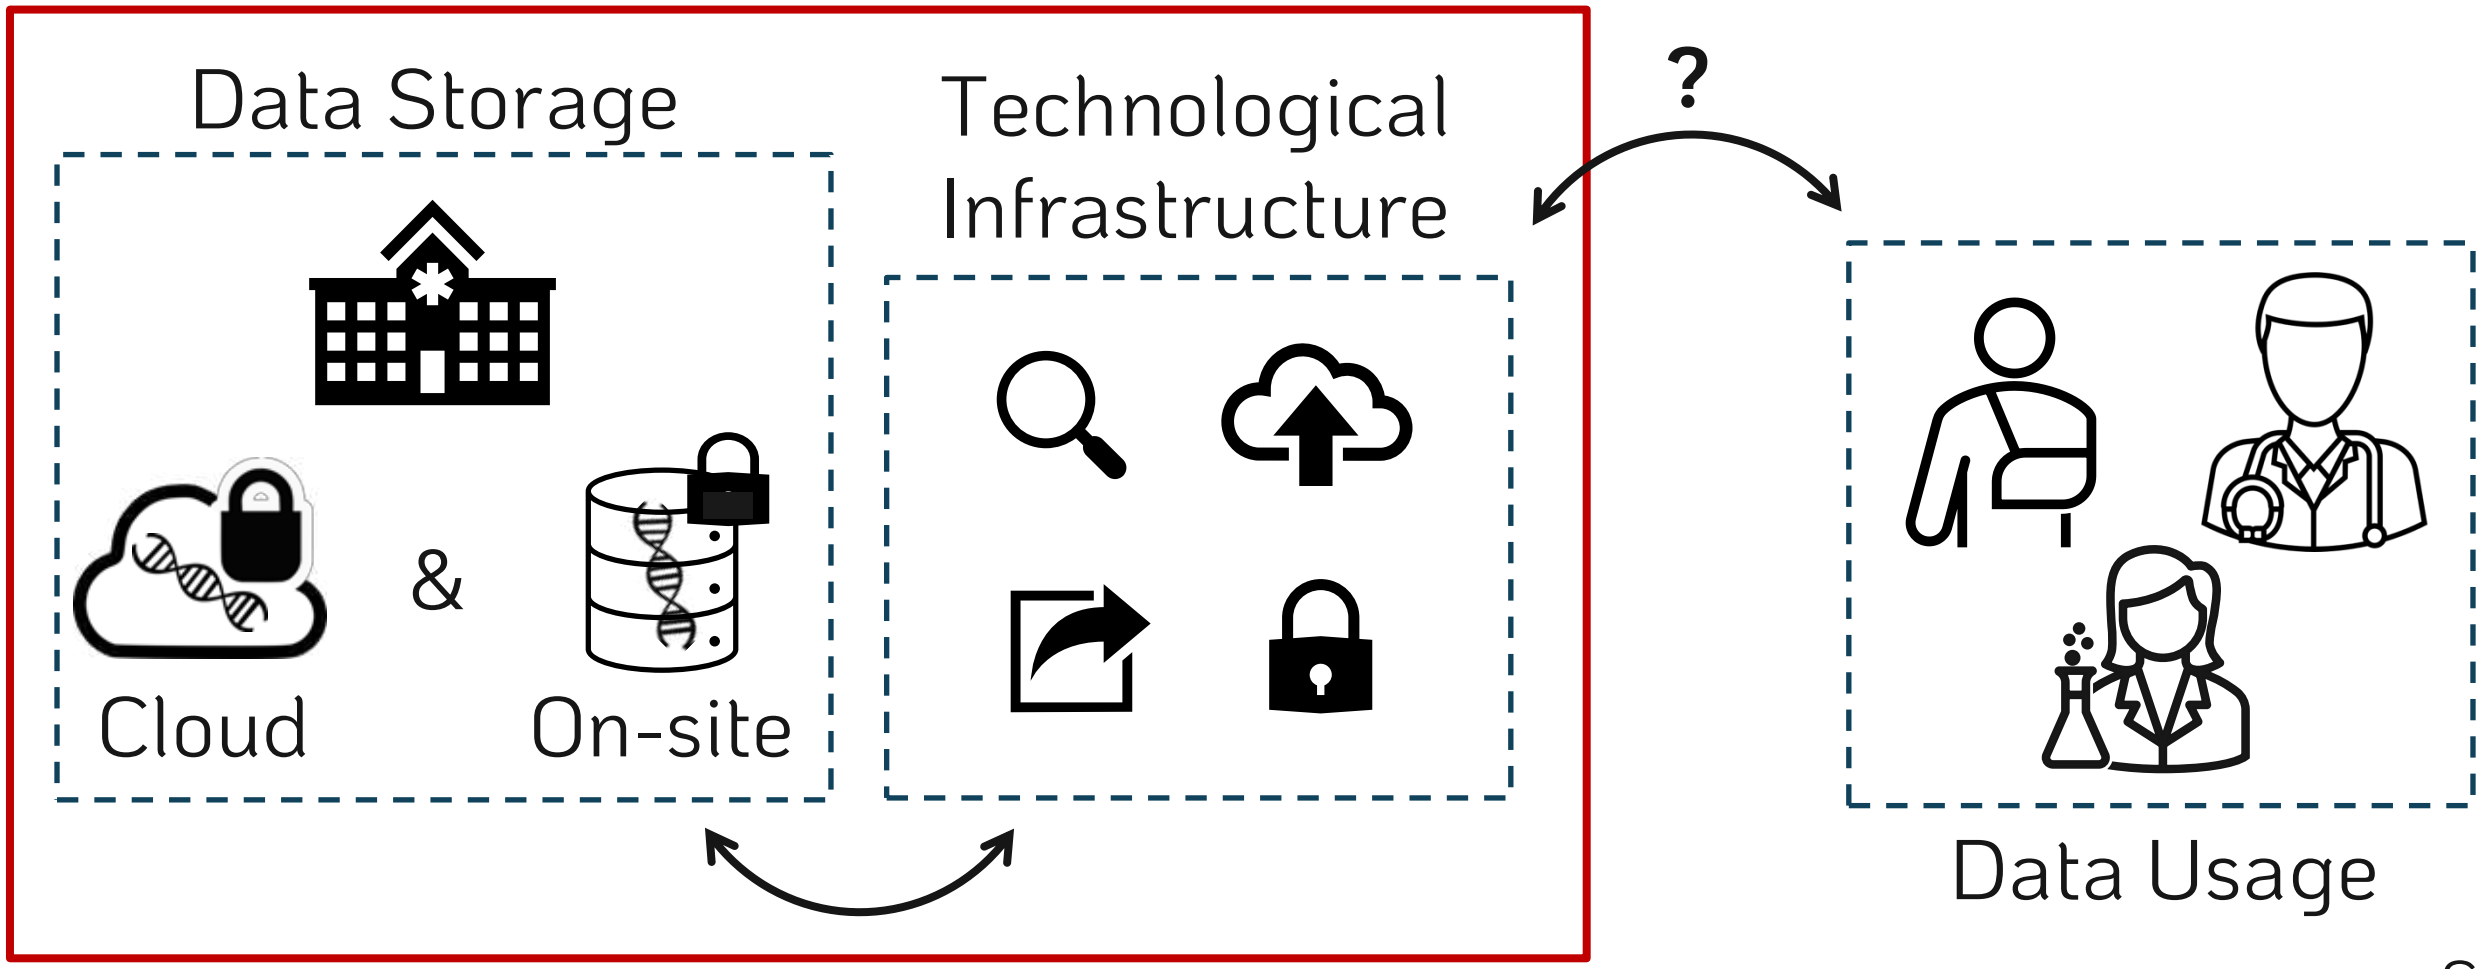

A [Hybrid] Solid Solution

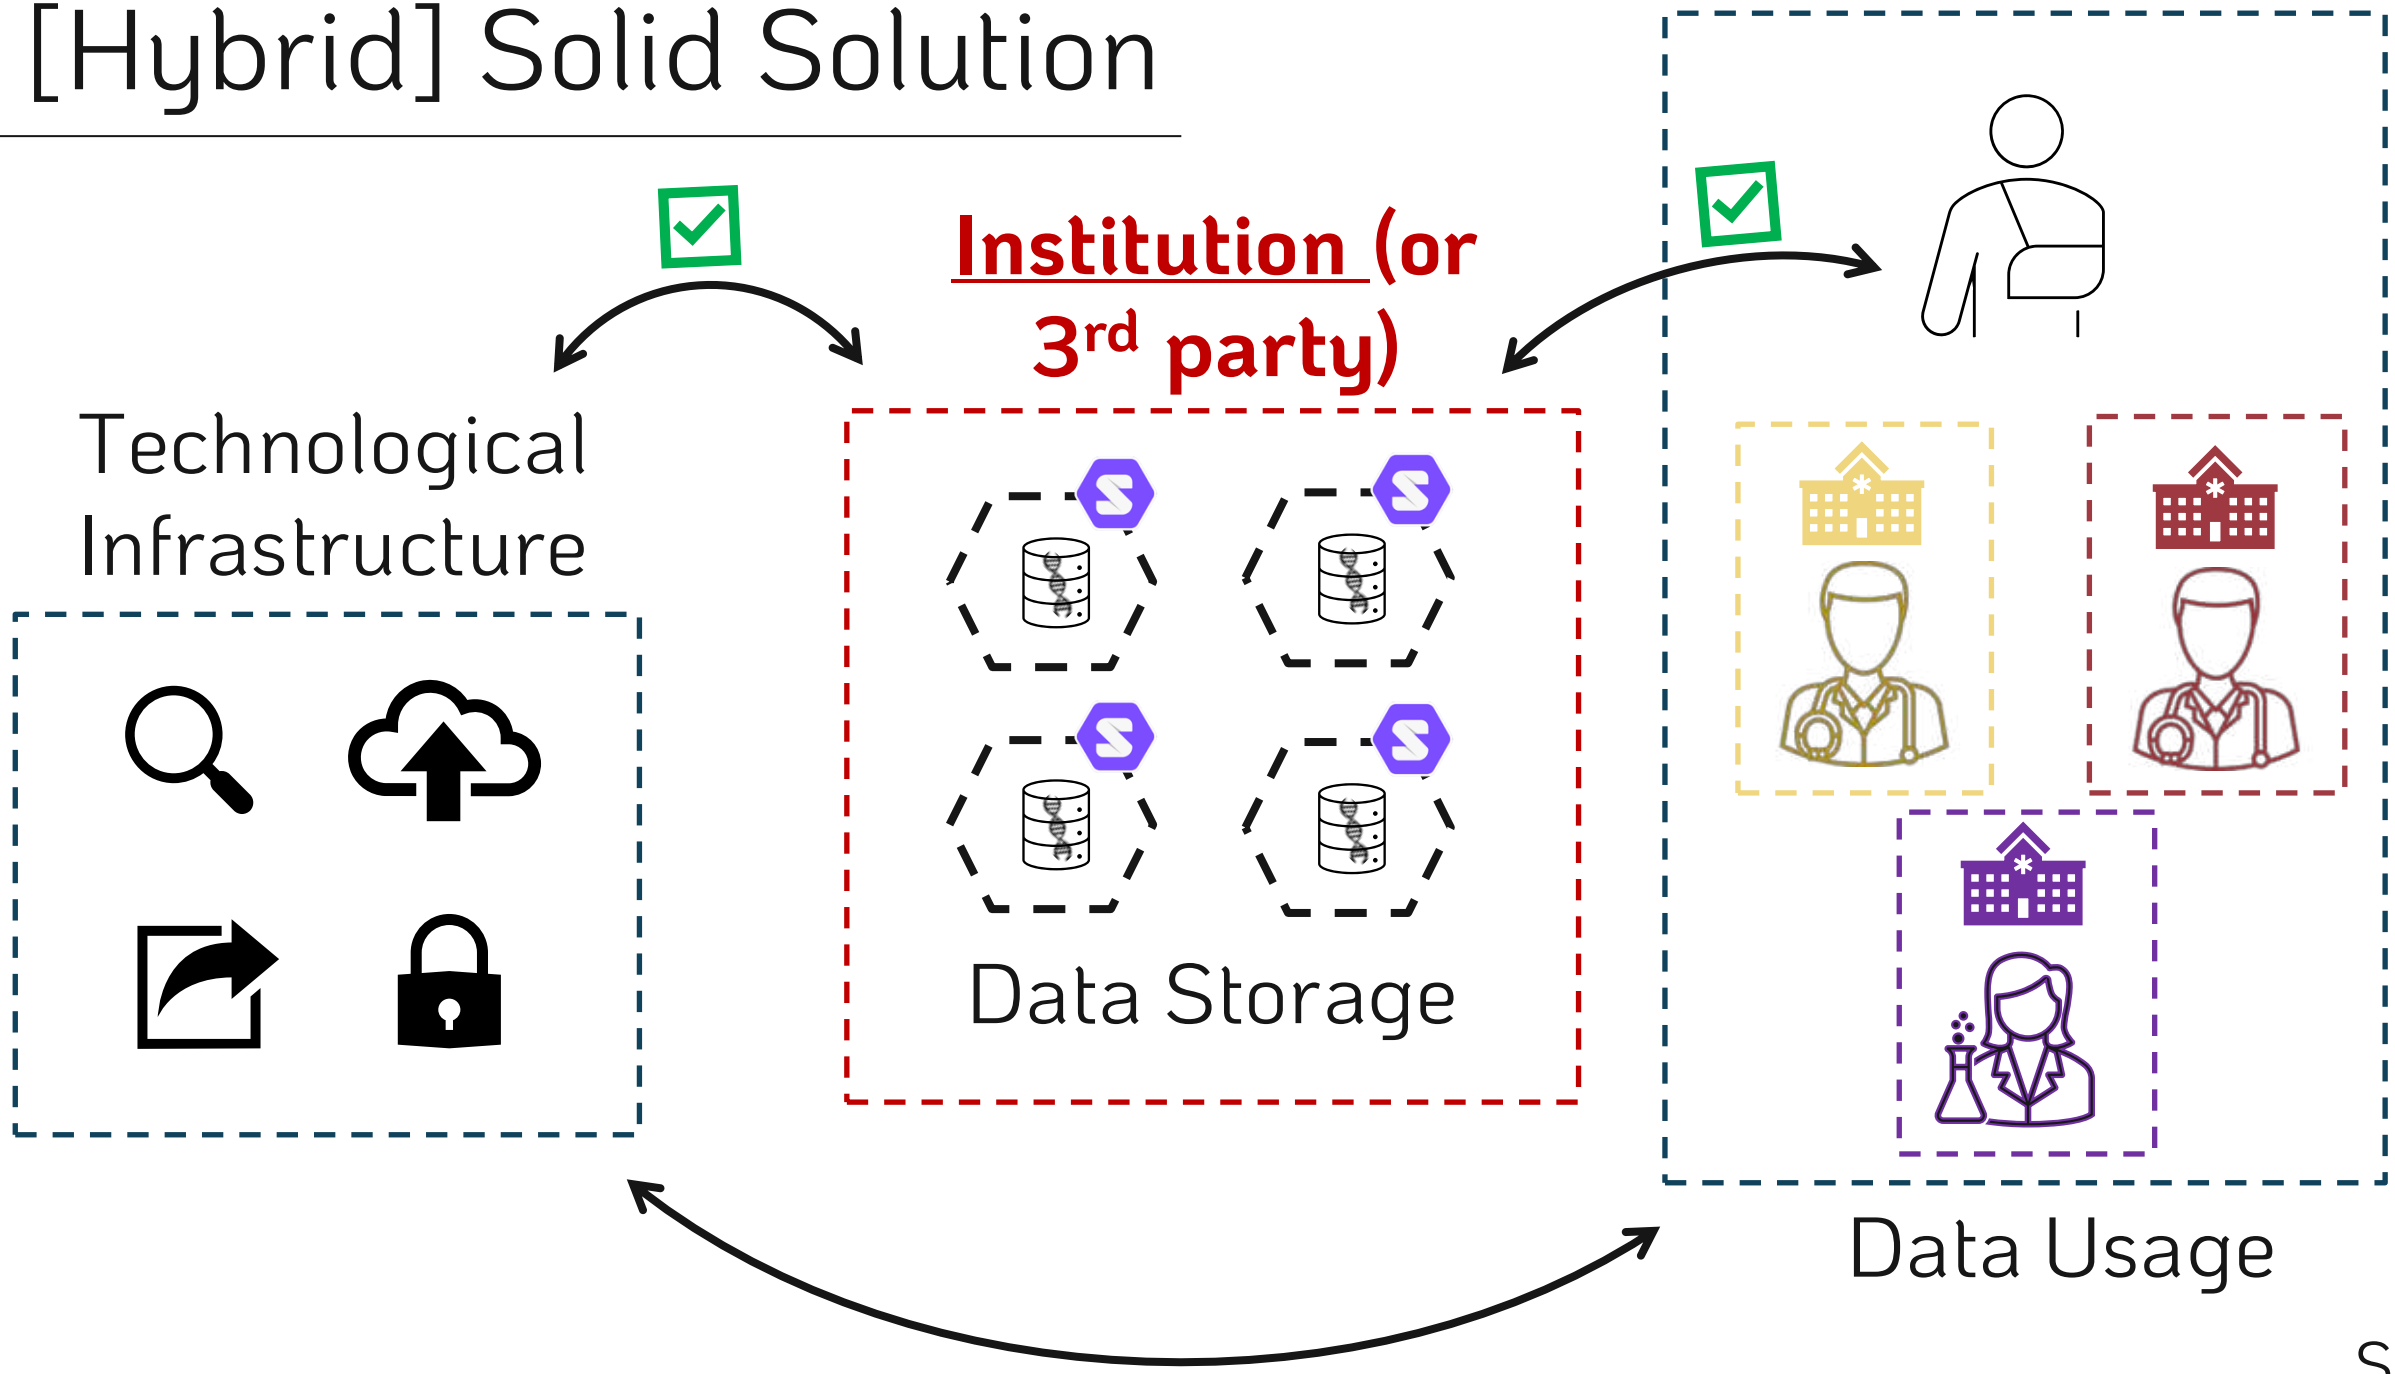

Solid Technology

- Does not solve ALL problems, but offers an environment where improvements are possible
  - ✓ Easier sharing (web-facilitated)
  - ✓ Granular privacy controls
  - ✓ Direct patient data access
  - ✓ Data representation flexibility (RDF)
  - ✓ Possibility of data-linking and querying

https://solidproject.org/TR/protocol

Project Goals

Create a proof-of-concept
storage **FRAMEWORK**
and accompanying **Web Application**

While **exploring representation + querying limits**

Envisioned Product Evolution

PhD End

Short Term

Medium Term

Long Term

PhD Deliverable

Start-up / Product development
+
Hospital implementation

Integrate with EHR
+
Increase application domains

Increase 3rd party application support
+
Scale to Europe and beyond

Workflow Overview

Data Storage

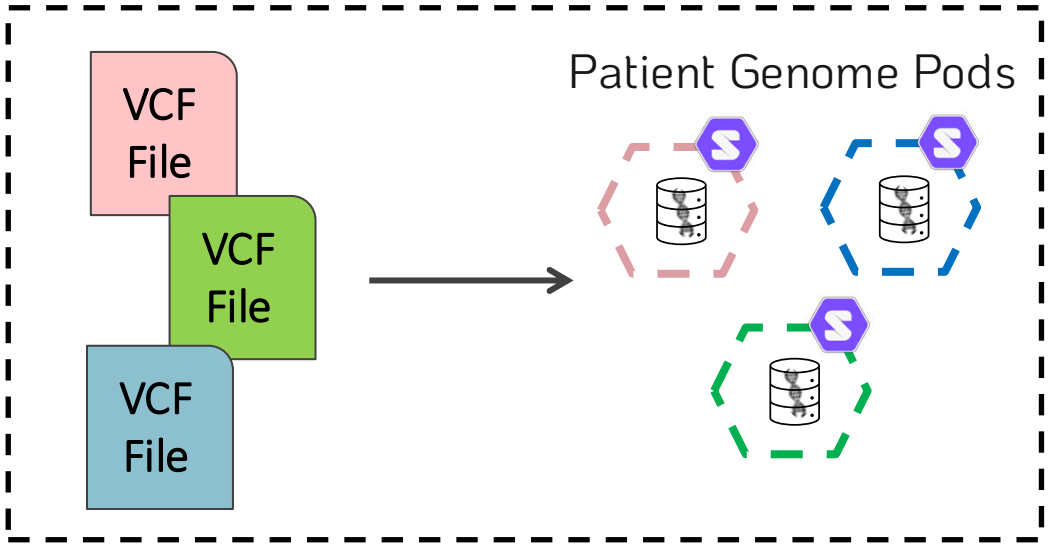

Data Storage

Solid can be implemented for genomic data storage.

Solid pod instances → Community Solid Server[1]

VCF Genomic Data → Illumina Platinum Genomes[2]

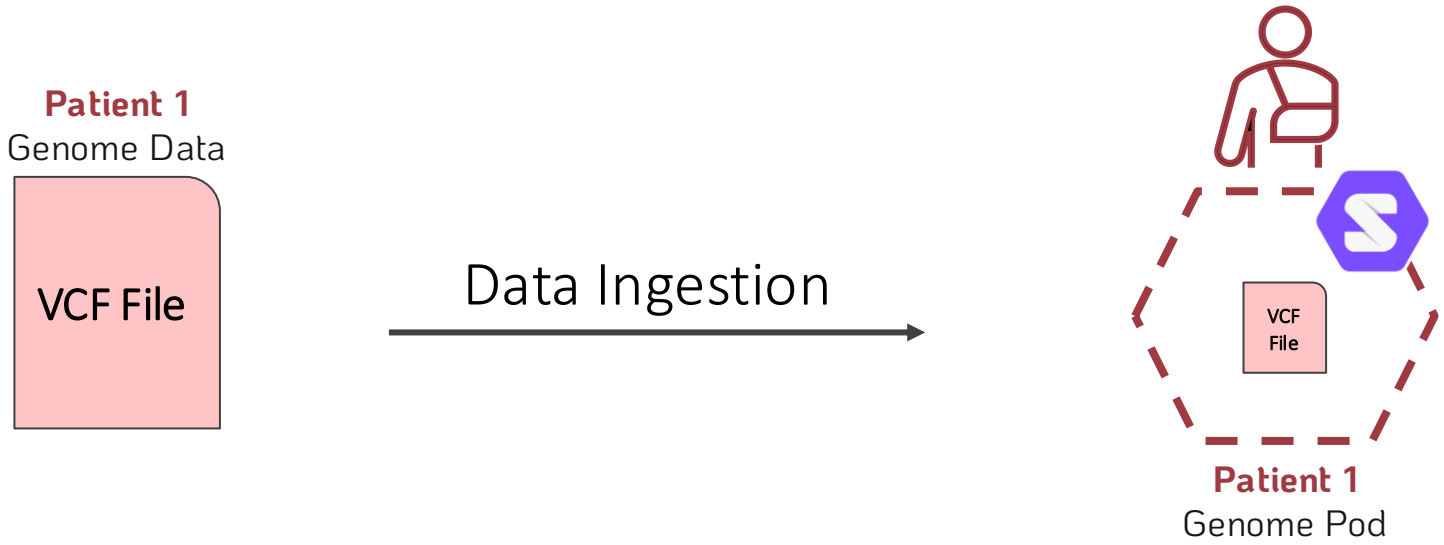

[1] https://github.com/CommunitySolidServer/CommunitySolidServer
[2] https://emea.illumina.com/platinumgenomes.html

Workflow Overview

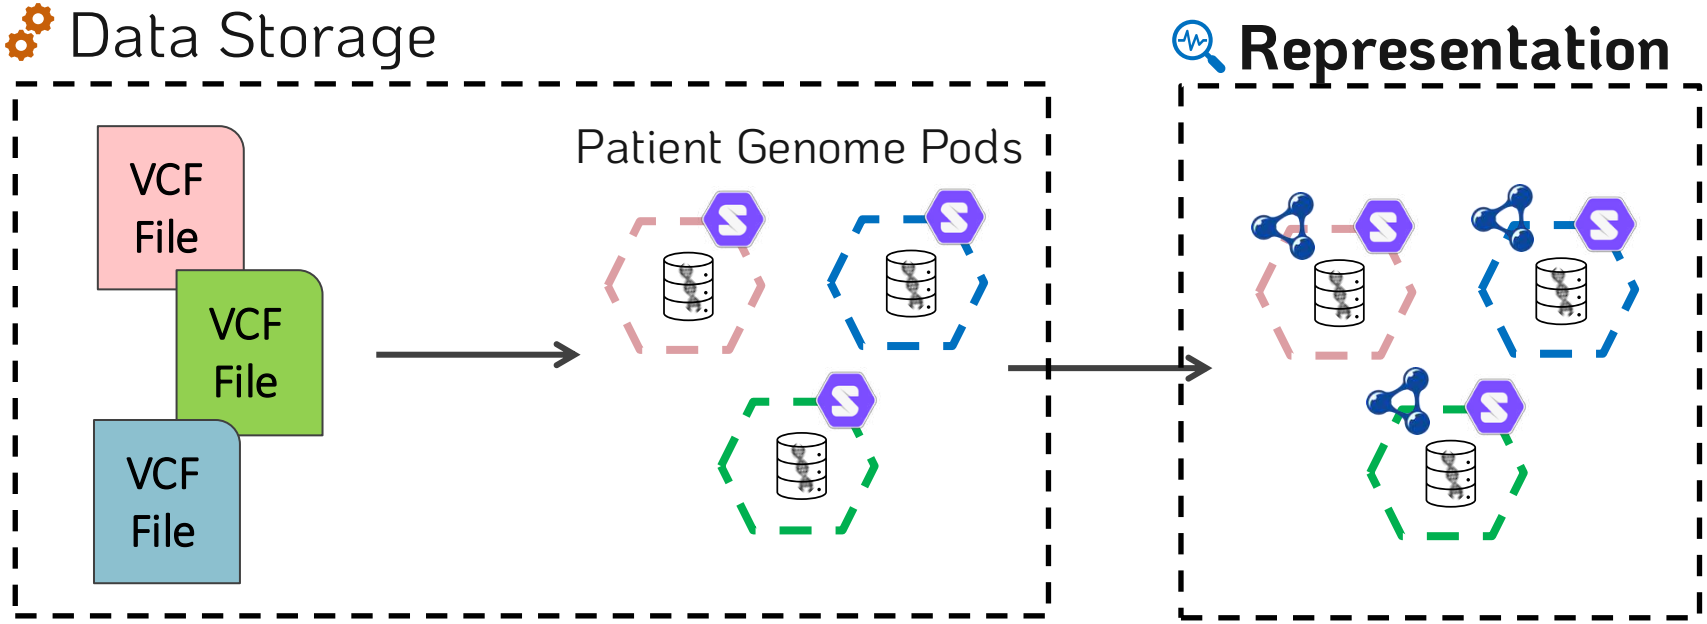

Genome Data in RDF

VCF genomic data can be represented as RDF.

VCF to RDF Vocabularies exist → SPHN Semantic Interoperability Framework[4]
+ FHIR HL7 Format[5]

VCF genomic data can be Linked to other data.

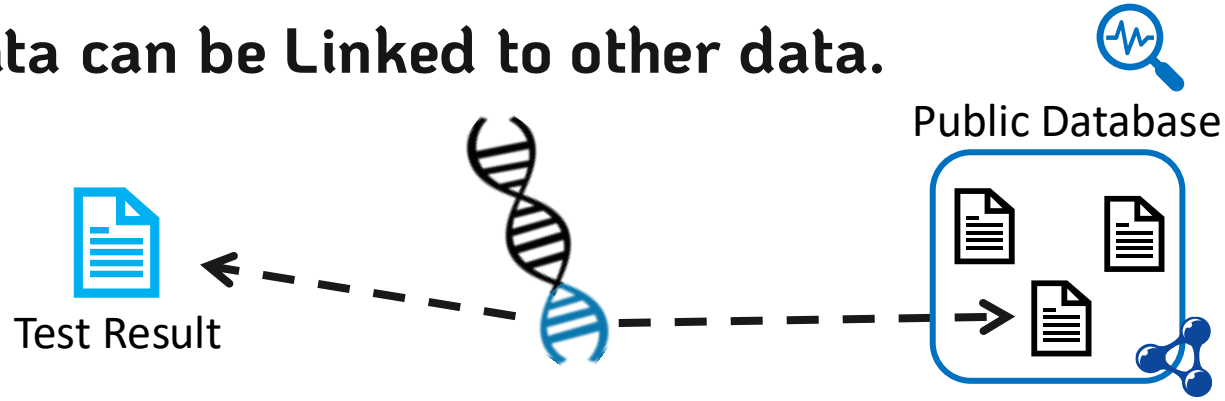

How to automate this?
What vocabularies could be used for the linkages?

[3] van der Horst E, et al. Bridging Clinical and Genomic Knowledge: An Extension of the SPHN RDF Schema for Seamless Integration and FAIRification of Omics Data. 2023.
[4] Dolin RH, et al. vcf2fhir: a utility to convert VCF files into HL7 FHIR format for genomics-EHR integration. 2021.

Workflow Overview

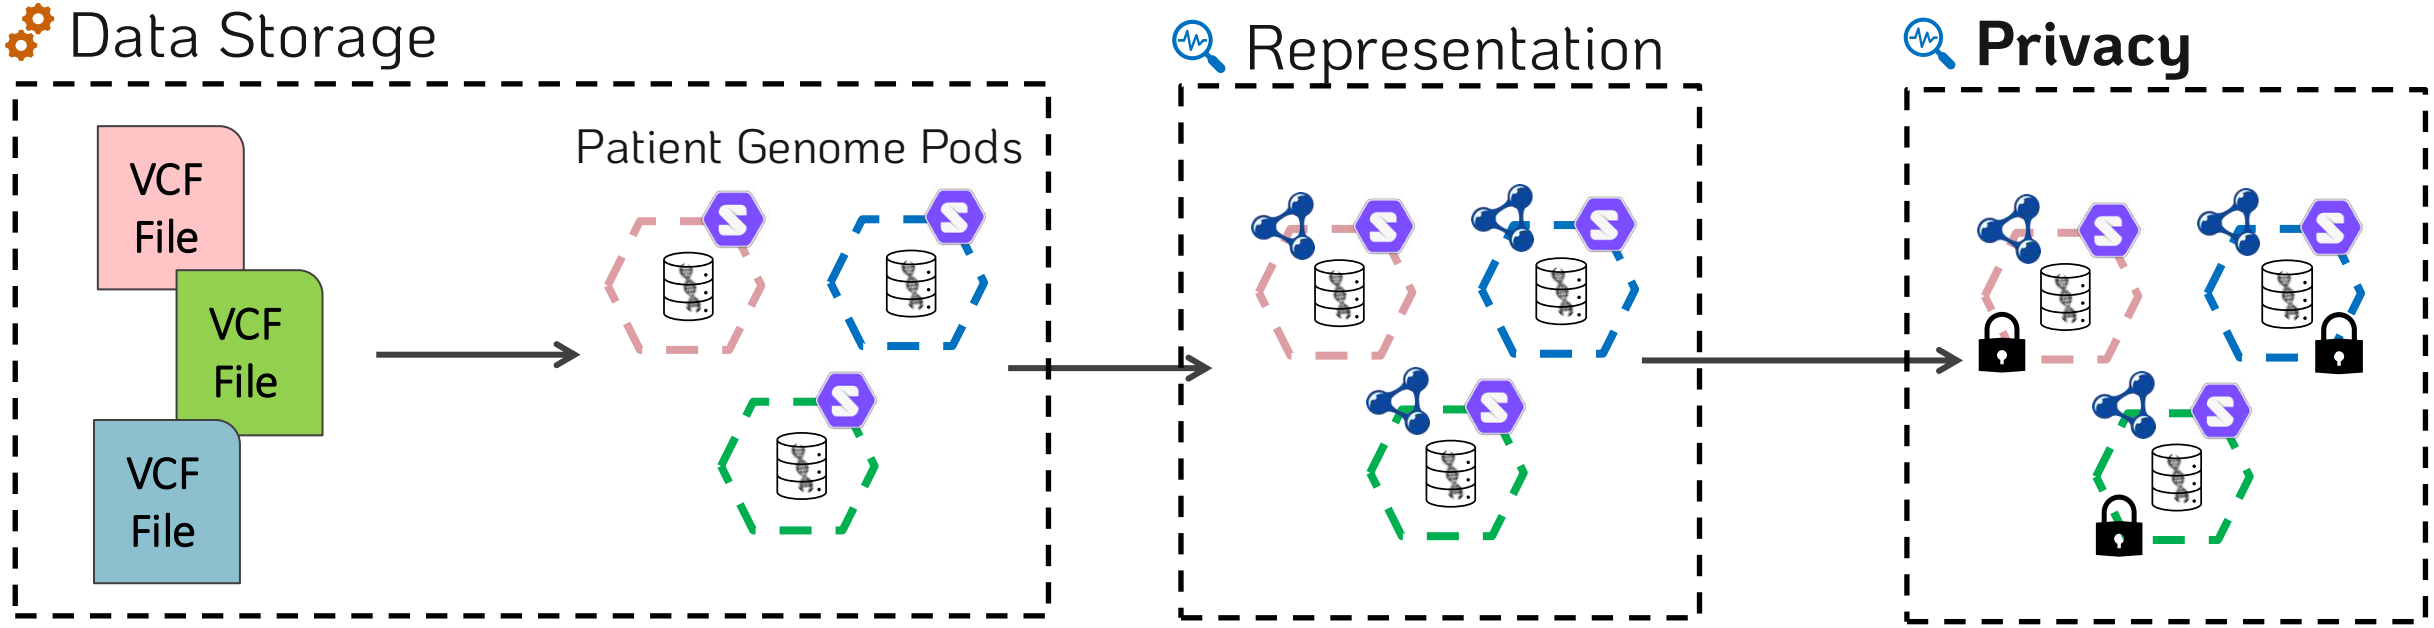

Data Privacy Controls

Granular privacy controls can be implemented in Solid for genomic data.

Web Access Contol → Solid Protocol ACP[3]

How granular can we make the privacy controls?

Privacy controls at chromosome level, or even mutation level.

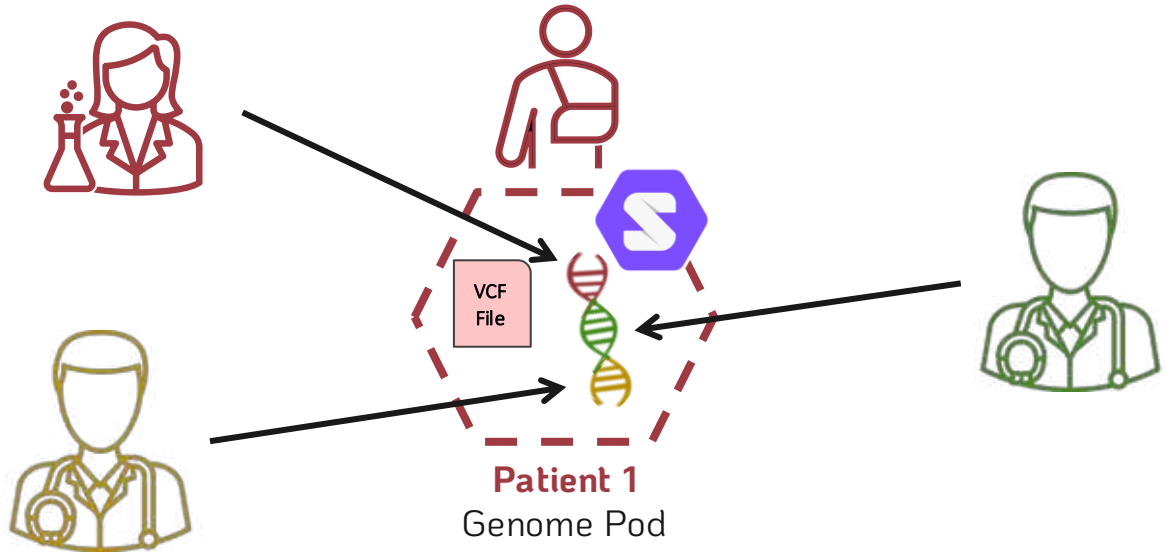

Chromosome is a semantic Class

Mutation is also a semantic Class, above a triple but below a Chromosome

[5] https://solidproject.org/TR/protocol

Querying Genome Pods

RDF genomic data pods can be queried.

Query Engine → Custom Comunica implementation/engine[6]

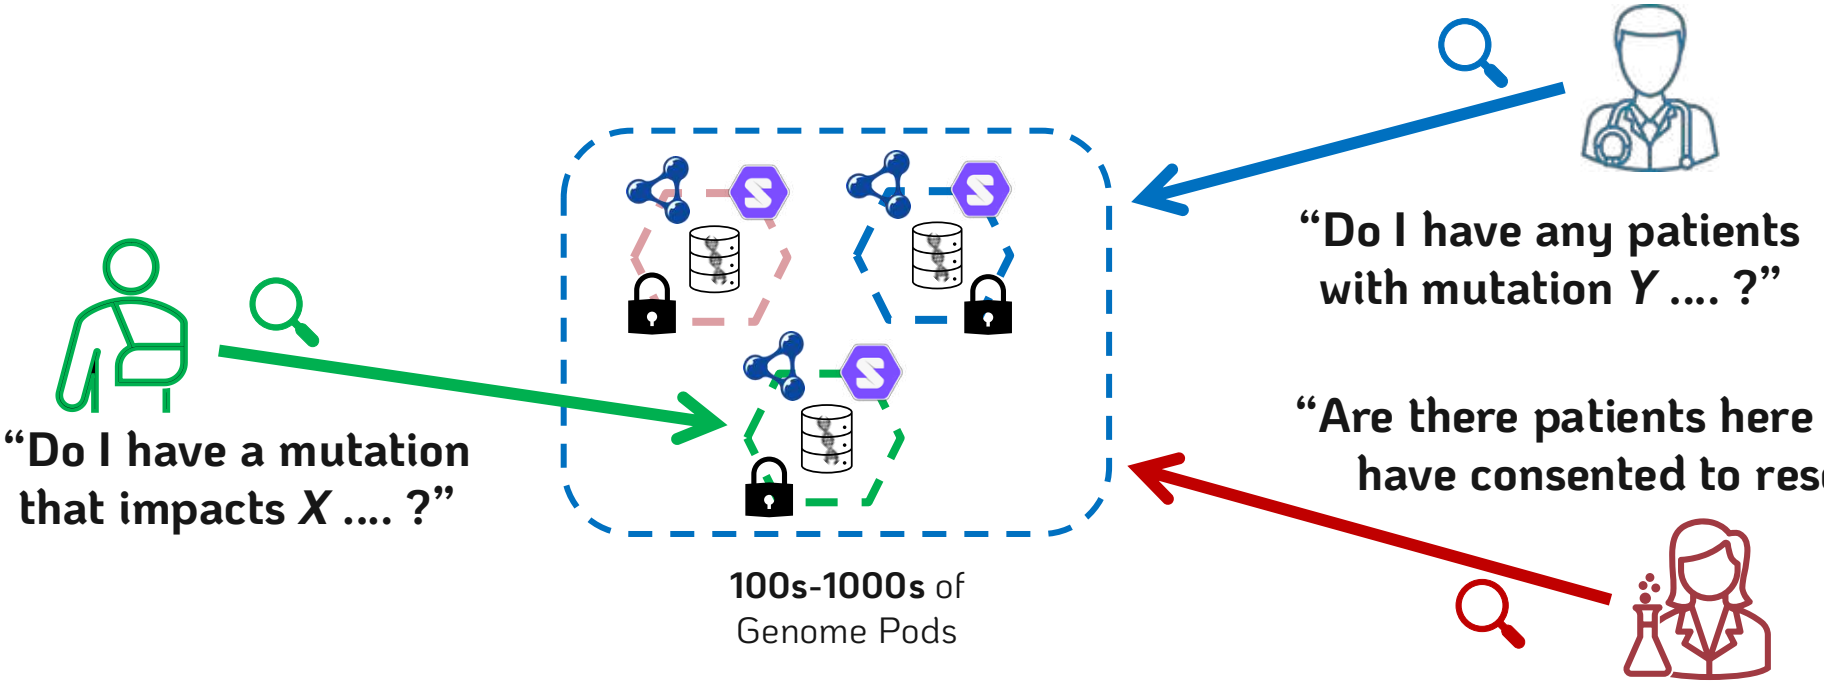

[6] https://comunica.dev/

Querying Genome Pods (possibilities)

Single Pod Query

Utilizes an intra-pod index

Many Pod Query

First, utilizes an extra-pod aggregator
Then, utilizes the intra-pod indexes

Beacon API for Researcher Queries

Query is translated from Beacon API to Comunica
Then, the **Many Pod** query approach is followed

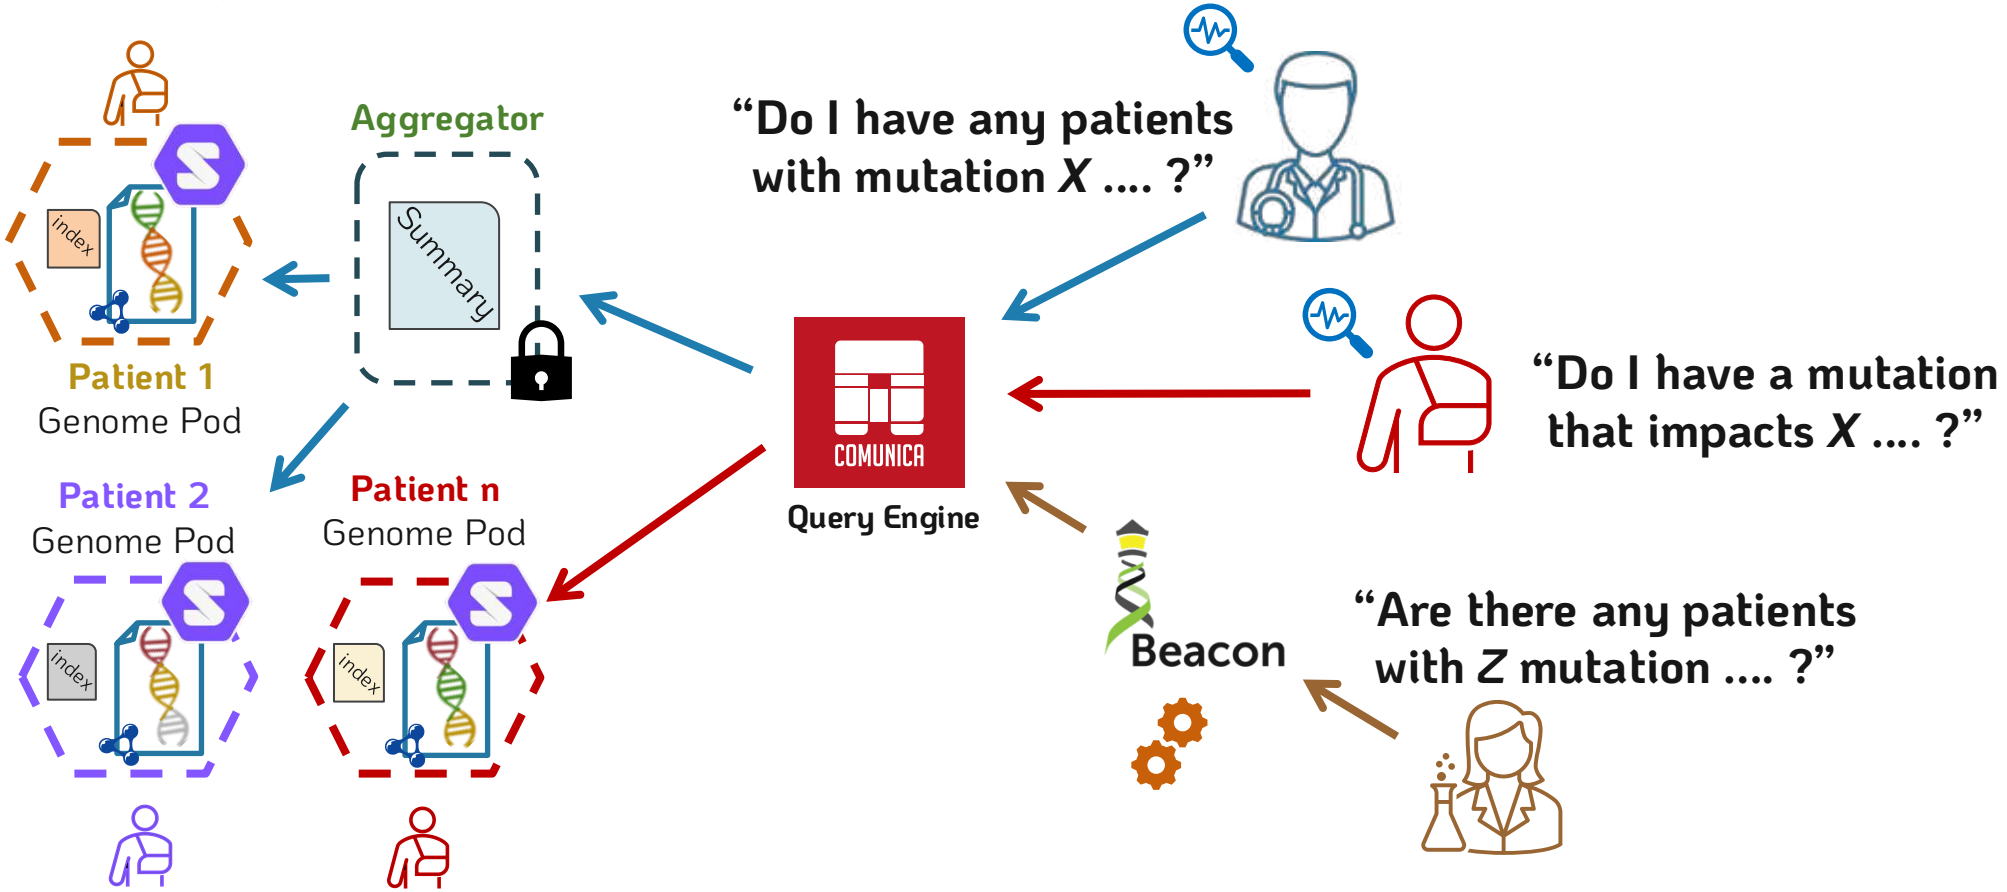

Implementation Specifics

Solid pod instances → Community Solid Server[1]

Genomic Data → Publicly available *VCF* data[2]

RDF Conversion → HERO Genomics + SPHN ontology[4]

Data-linking → Actively exploring ontologies / methodologies

Data Querying → Custom instance of Comunica[6]